


Transformers for Medical Imaging

Chengzhi Shen

MS.c. Biomedical Computing



Outline

- Self-attention & Transformers[1]
- Vision Transformers[2]
- Applying Transformers in Medical Imaging[3][4]

[1] *Attention is all you need*, A. Vaswani, et al., *Advances in Neural Information Processing Systems* pp. 5998-6008, 2017

[2] *An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale*, A. Dosovitskiy, et al. *arXiv:2010.11929*, 2020

[3] *UNETR: Transformers for 3D Medical Image Segmentation*, A. Hatamizadeh et al., *arXiv:2103.10504v1*, 2021

[4] Shamshad, Fahad, et al. "Transformers in medical imaging: A survey." *arXiv preprint arXiv:2201.09873* (2022).

Attention is All You Need

Attention is All You Need

- First Transformer model, designed for text sequences in Natural Language Processing (NLP)
- Based on **self-attention**
- Considered as **fundamental** AI architectures like CNN, RNN
- A large number of following up works:
 - Language models: BERT, GPT etc.
 - Other modalities: VisionTransformers, Graphormers...



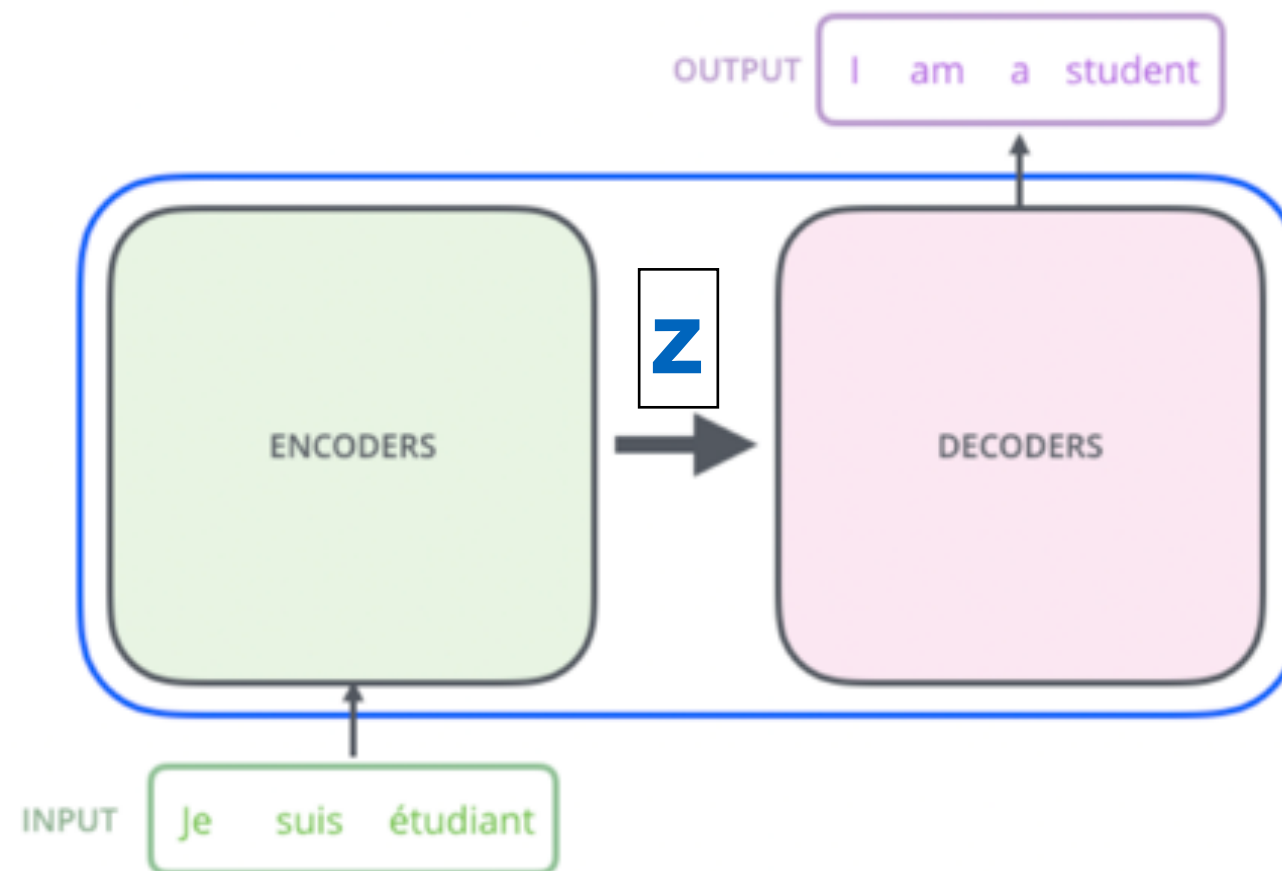
Motivation

- The **cat** drank the milk because **it** was hungry.
- The cat drank the **milk** because **it** was sweet.



Background: seq2seq models

- The encoder-decoder architecture
- **Encoder**: input sequence $X = (x_1, \dots, x_n)$, get representation z
- **Decoder**: given z , generate output sequence $Y = (y_1, \dots, y_m)$



Transformer Architecture

- Encoder - Decoder

N attention blocks

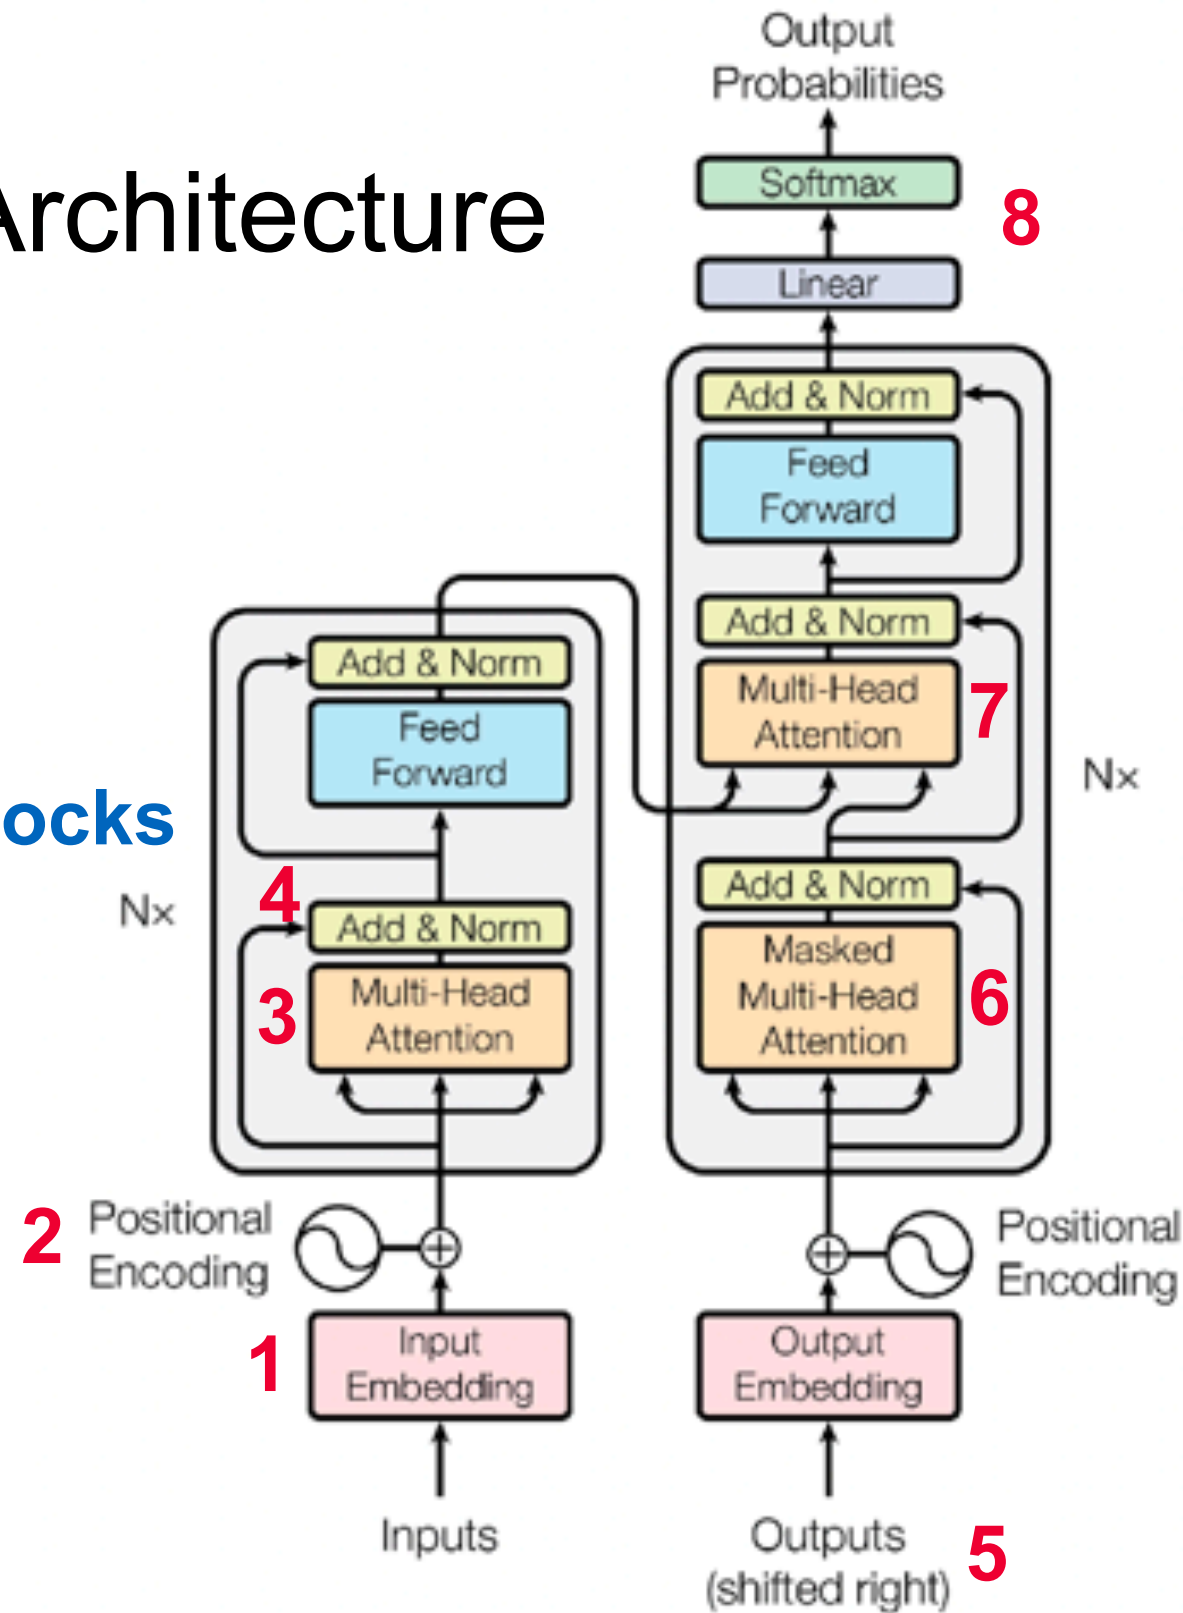


Figure 1: The Transformer - model architecture.

Transformer Architecture

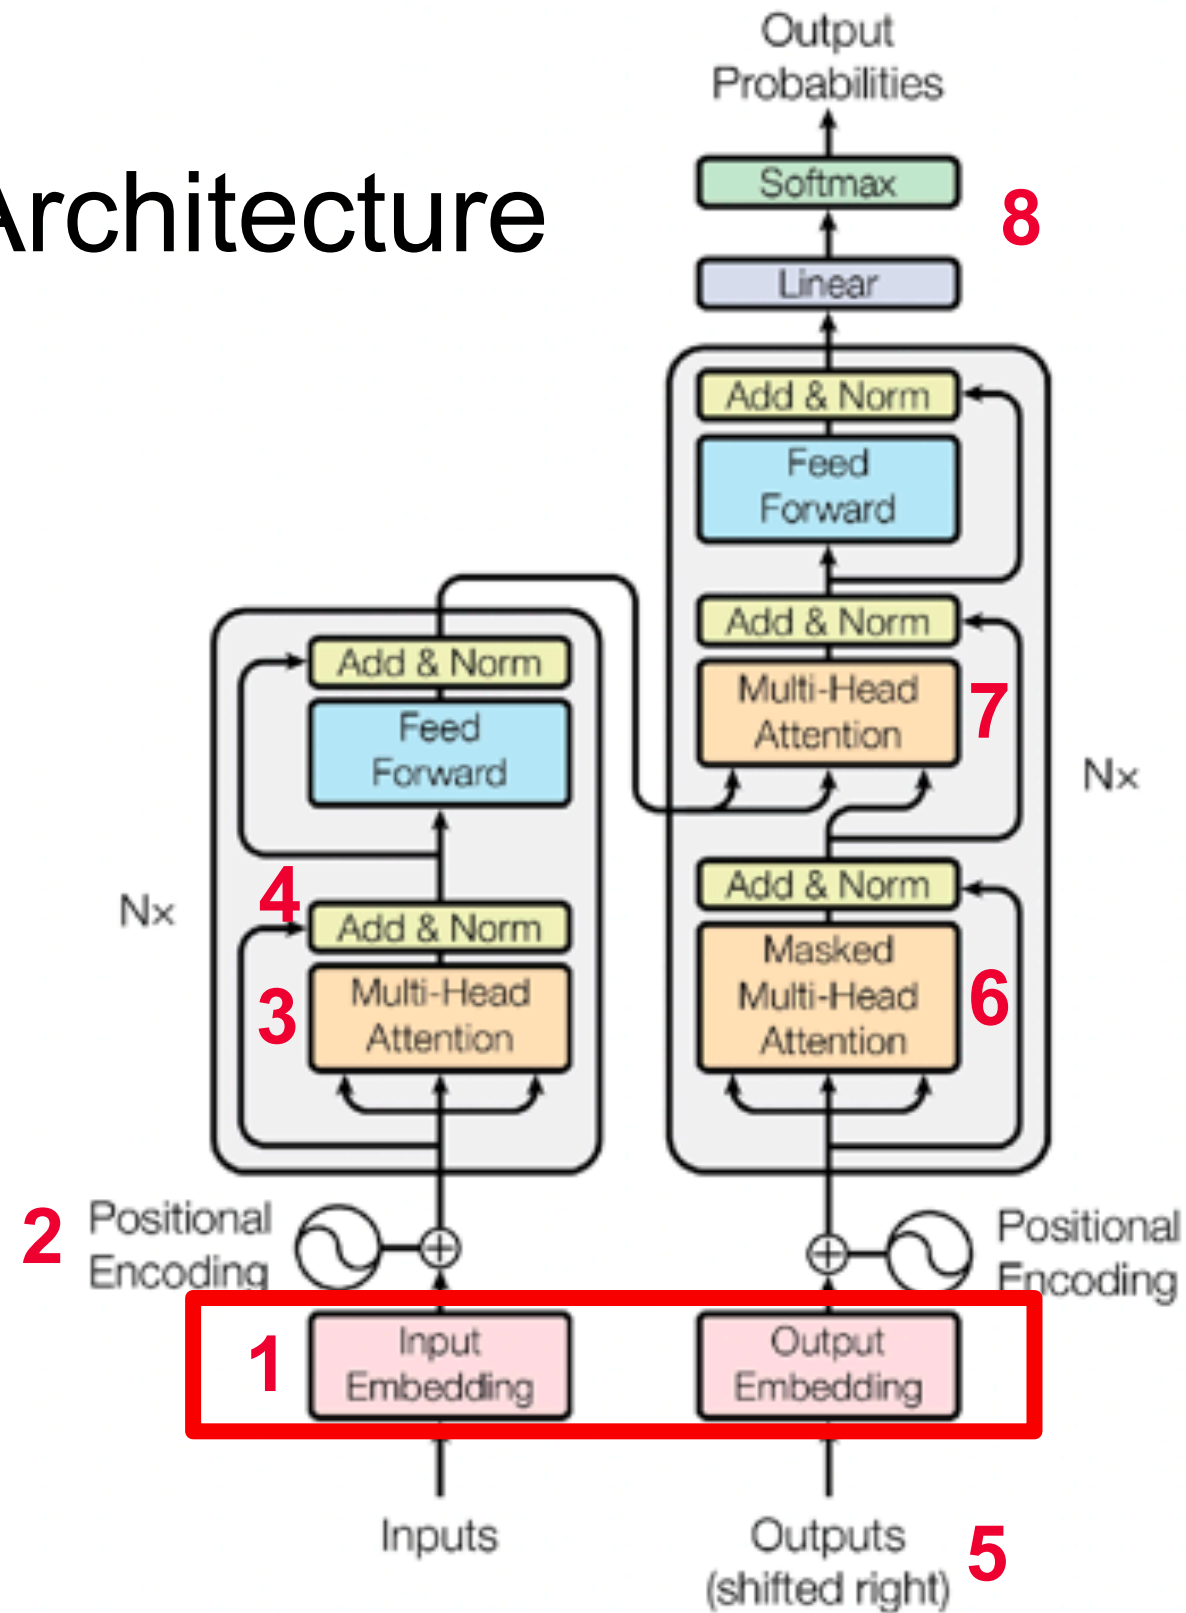
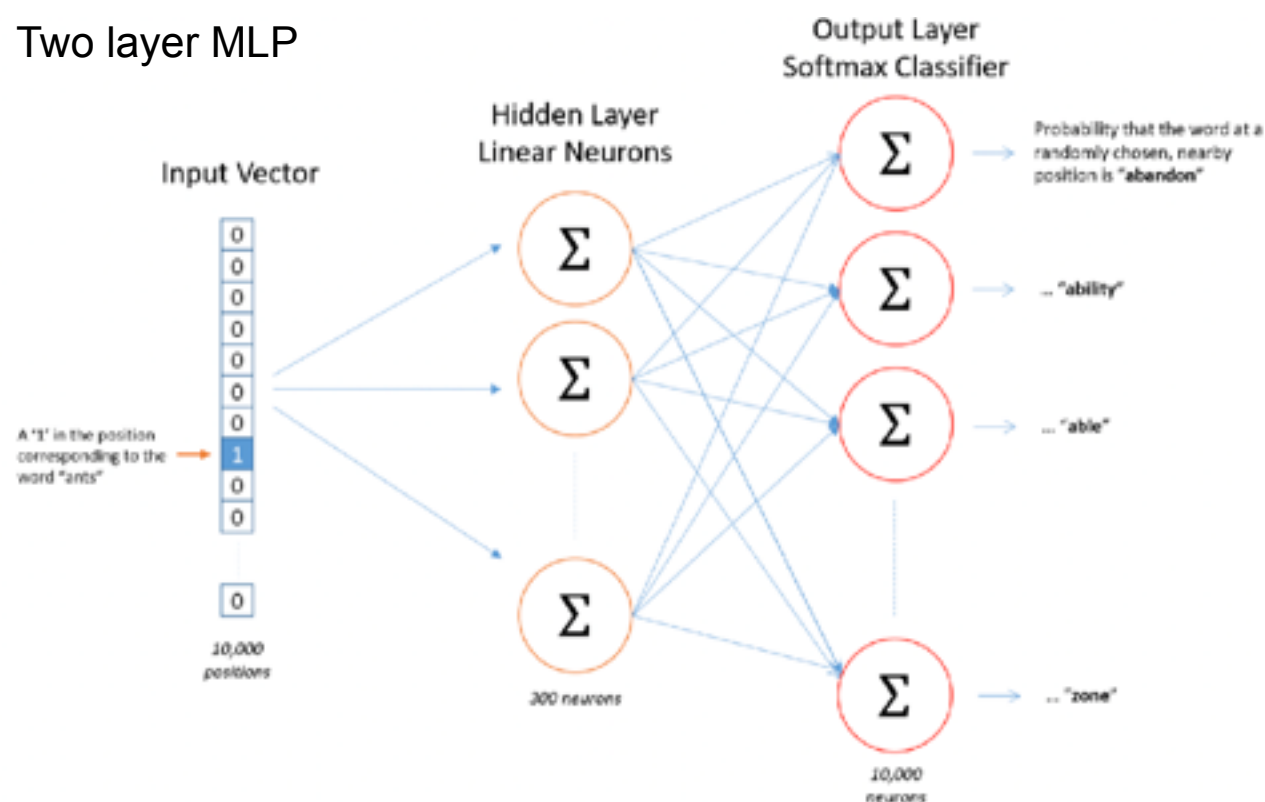


Figure 1: The Transformer - model architecture.

Embeddings

- Intuition: convert texts into vectors(embeddings) that can be handled by computer
- One-hot encoding fails: too sparse when having large vocabularies
- Learnable embeddings: Word2vec
 - CBOW, skip-gram


Two layer MLP



The embedding matrix

$$\text{Word position} \begin{bmatrix} 0 & 0 & 0 & 1 & 0 \end{bmatrix} \times \begin{bmatrix} 17 & 24 & 1 \\ 23 & 5 & 7 \\ 4 & 6 & 13 \\ 10 & 12 & 19 \\ 11 & 18 & 25 \end{bmatrix} = \text{Word embedding} \begin{bmatrix} 10 & 12 & 19 \end{bmatrix}$$

The more similar, the closer



Transformer Architecture

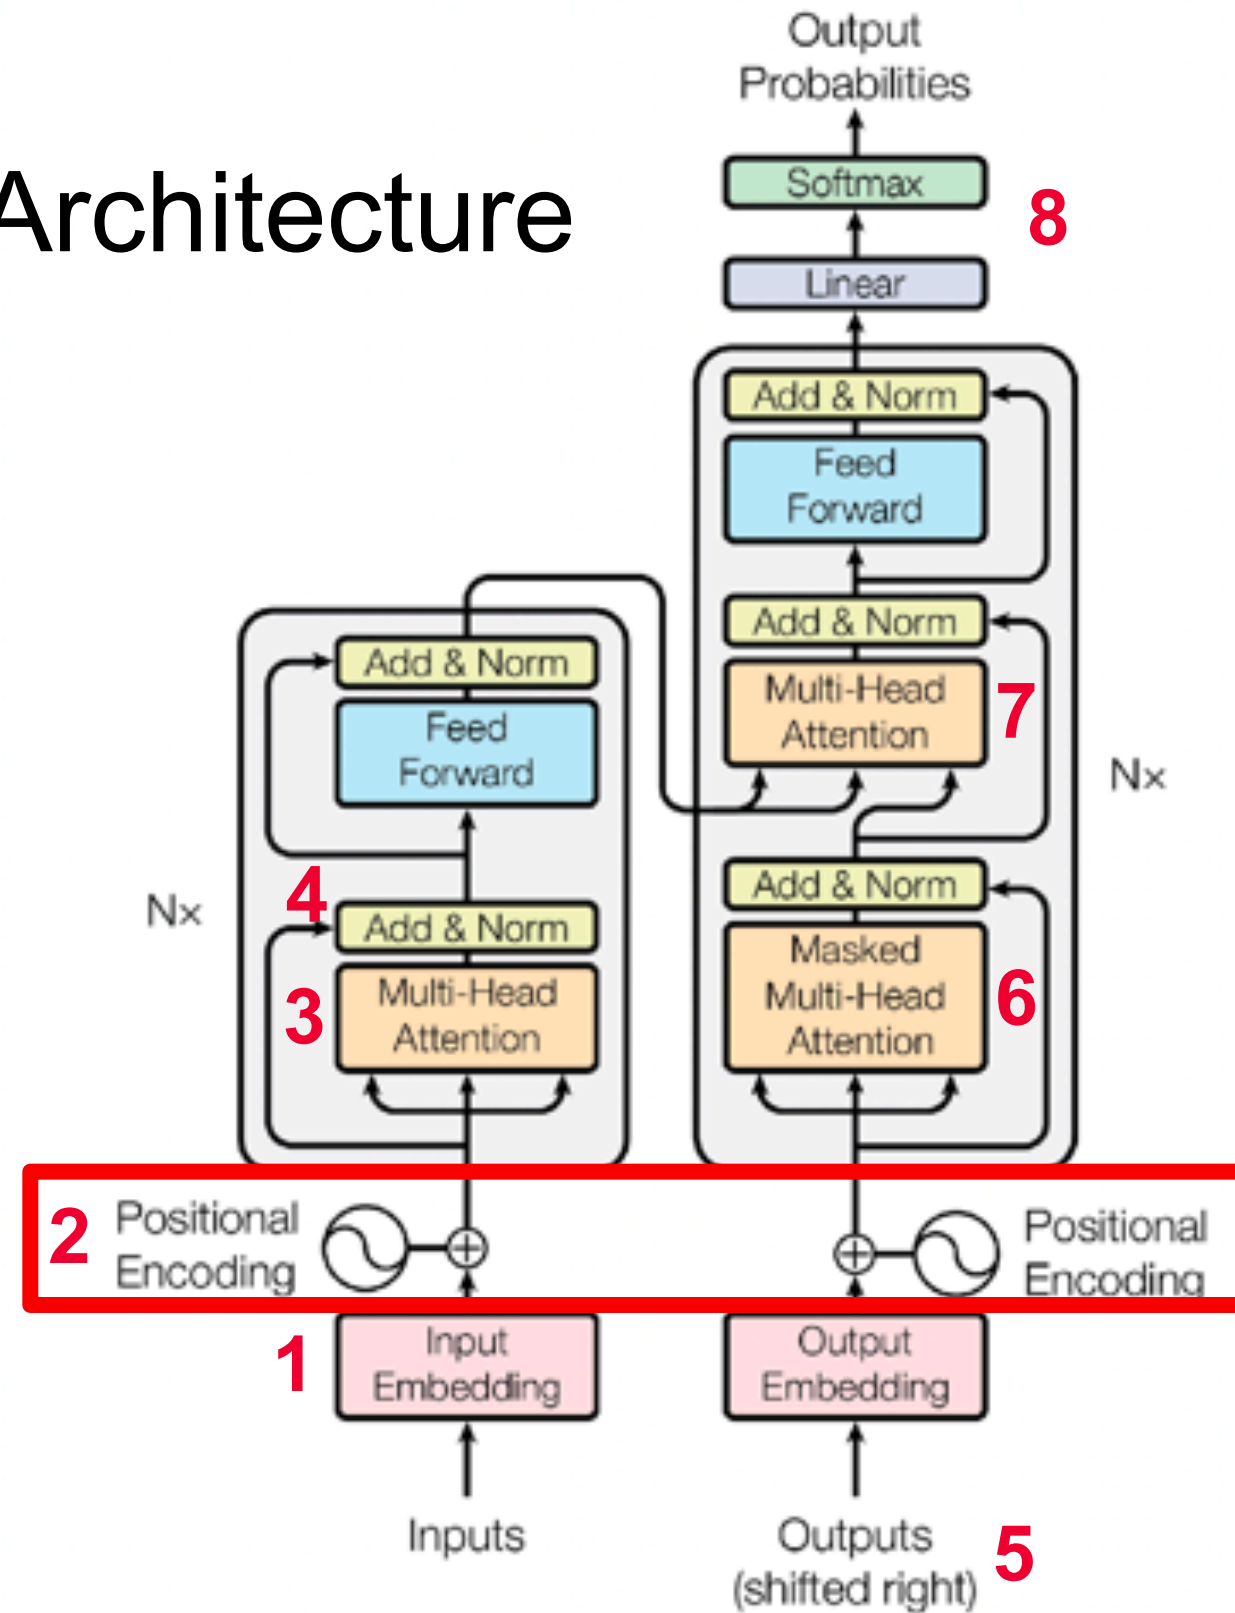



Figure 1: The Transformer - model architecture.

Positional Encodings

- **Order of words** matter in the sentence!!
- Same dimension as input embeddings, directly sum up with them



*“Since our model contains no recurrence and no convolution, in order for the model to make use of **the order of the sequence**, we must inject some information about the **relative or absolute position** of the tokens in the sequence.”*


Positional Encodings

- Sin and Cos functions:

$$PE_{(pos, 2i)} = \sin(pos/10000^{2i/d_{\text{model}}})$$

$$PE_{(pos, 2i+1)} = \cos(pos/10000^{2i/d_{\text{model}}})$$

where pos is the position and i is the dimension. **Fixed values that depend only on the max length of the sequence.**



Orange: 0th word's encoding for a 40-word sequence.
Blue: 1st word's encoding for a 40-word sequence.

- **Also can be learnable (later in VisionTransformer...)**

Transformer Architecture

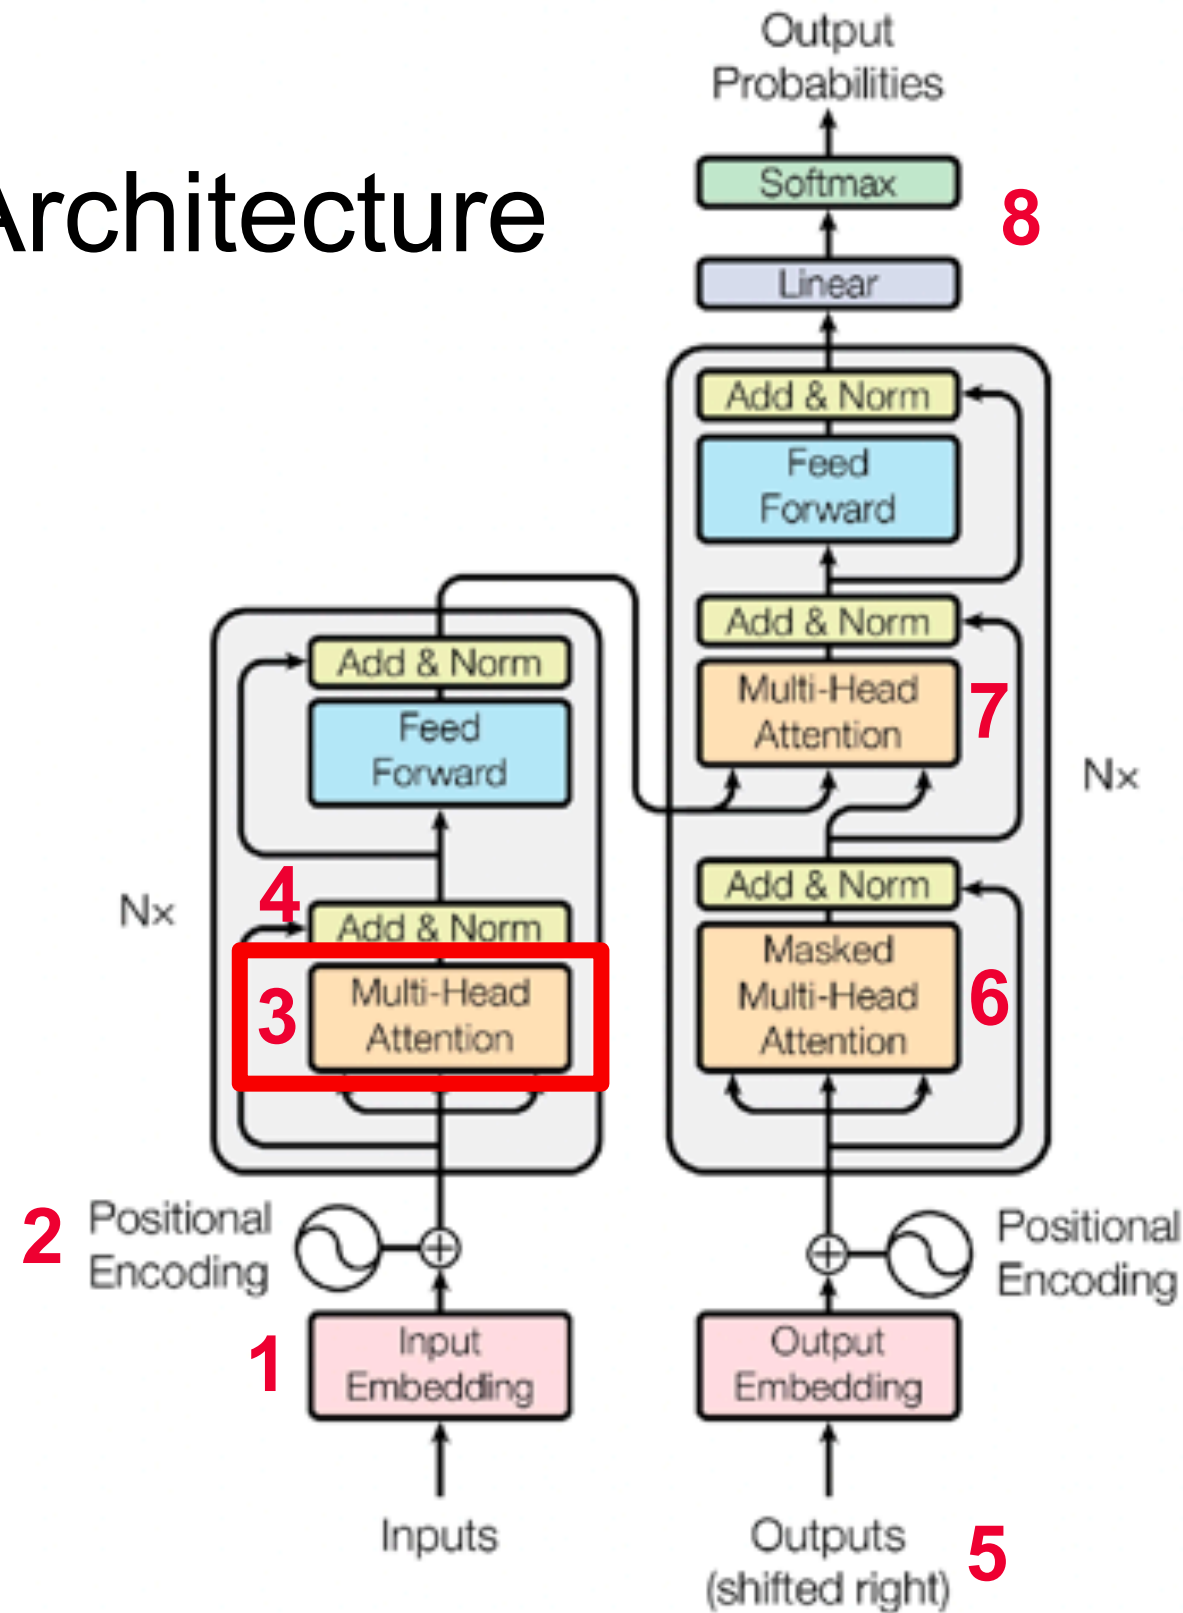


Figure 1: The Transformer - model architecture.

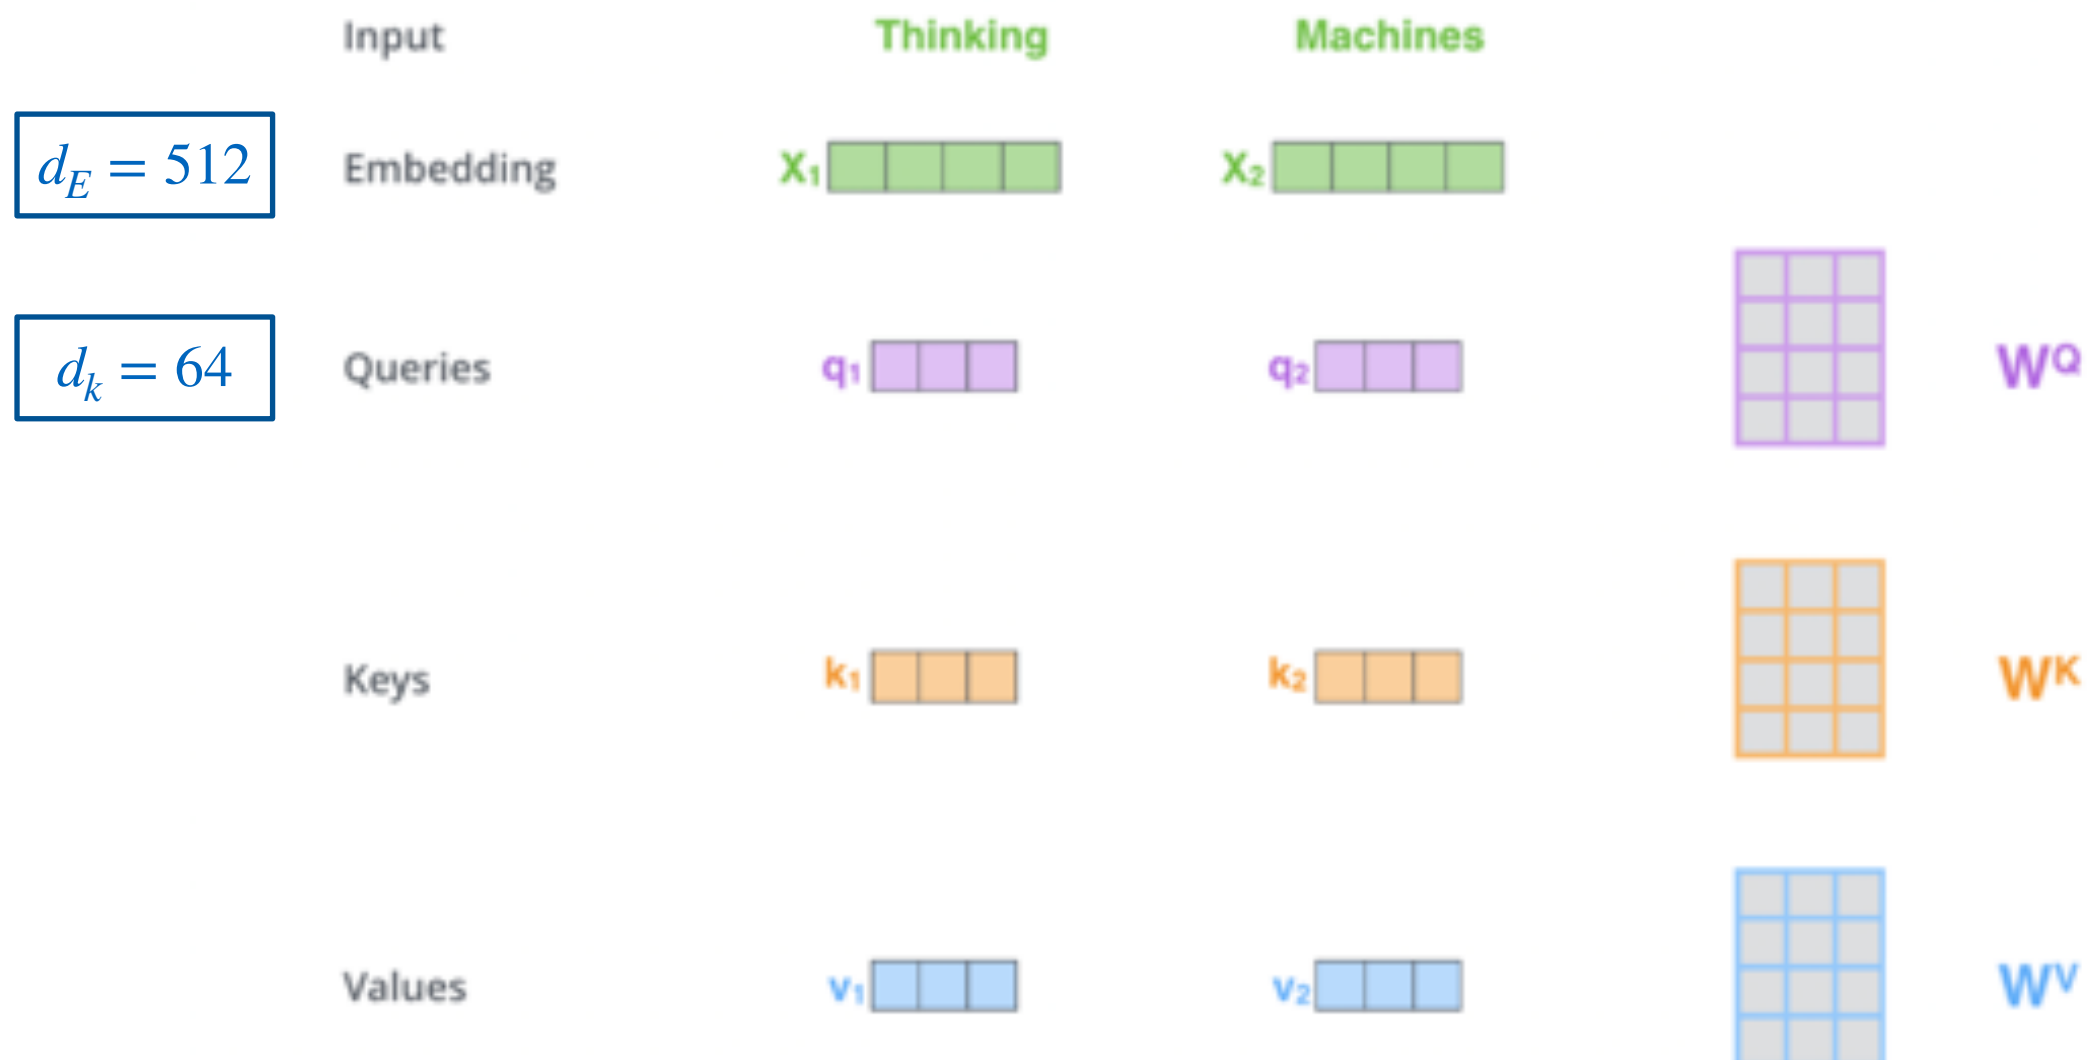
Self-Attention

- Intuition: to understand the meaning of the current word, we should look at other words in the sentence
 - The **cat** drank the milk because **it** was hungry.
 - The cat drank the **milk** because **it** was sweet.

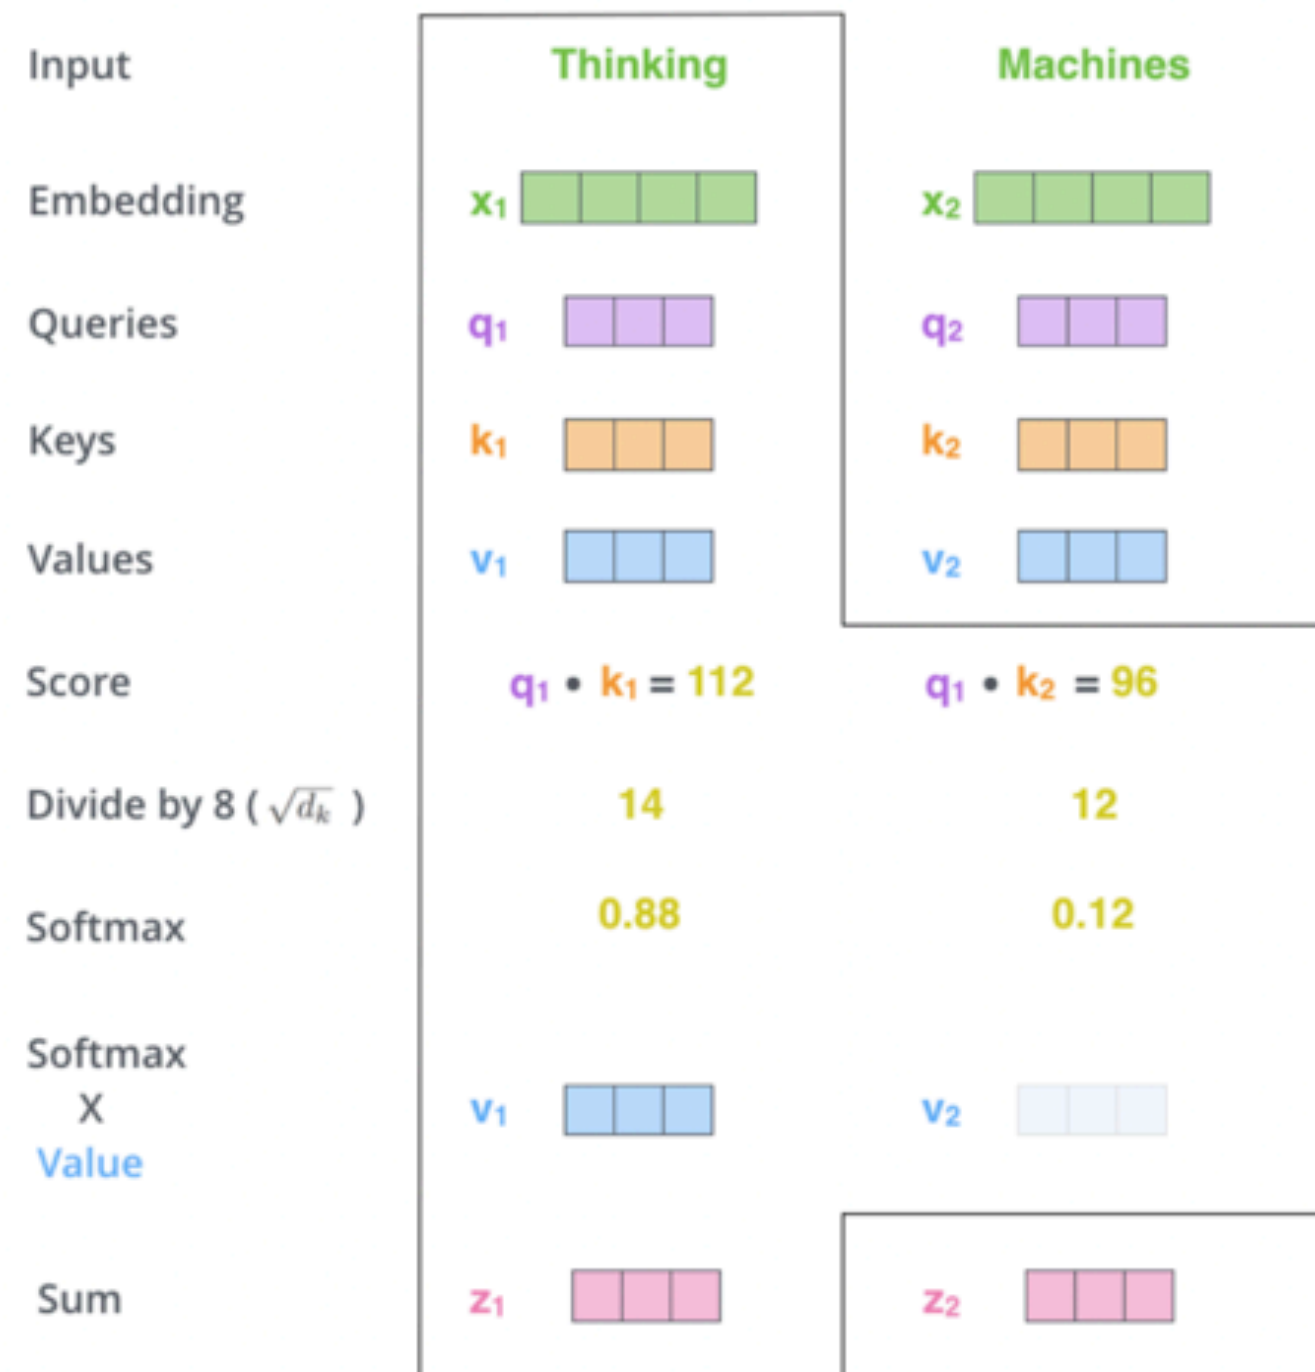


Self-Attention

- Initialize a **Query** W^Q , a **Key** W^K , and a **Value** W^V
- Matrix multiplication with word embeddings \rightarrow reduce dimensionality




Self-Attention



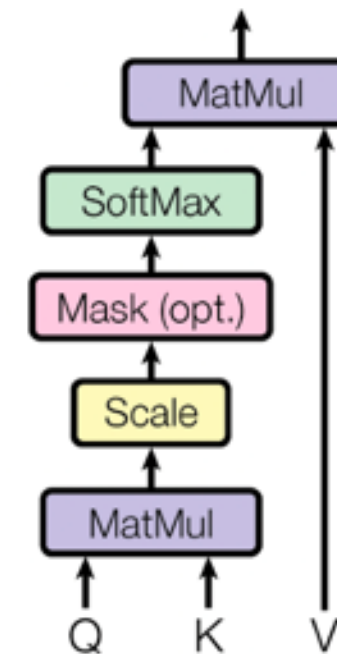
Self-Attention

- General formula and matrix calculations:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V$$



Scaled Dot-Product Attention



Multi-Head Self-Attention

1) This is our input sentence*




2) We embed each word*

3) Split into 8 heads. We multiply X or R with weight matrices

4) Calculate attention using the resulting $Q/K/V$ matrices

5) Concatenate the resulting Z matrices, then multiply with weight matrix W^O to produce the output of the layer

Thinking
Machines



...

...

...

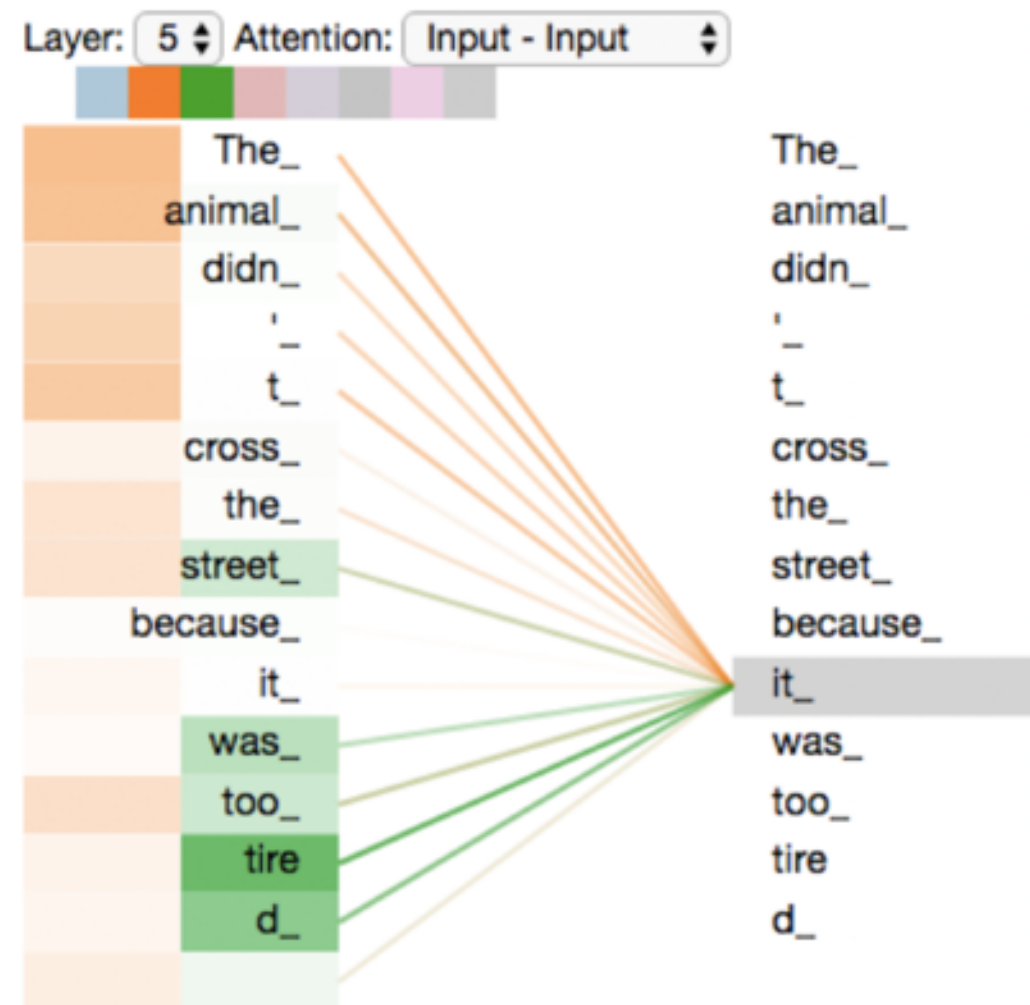


Multi-Head Self-Attention

- Intuition: different heads care for different aspects of the sentence.

*“The animal didn’t cross the street because **it** was too tired.”*

- Head0: the animal
- Head1: tired



Transformer architecture

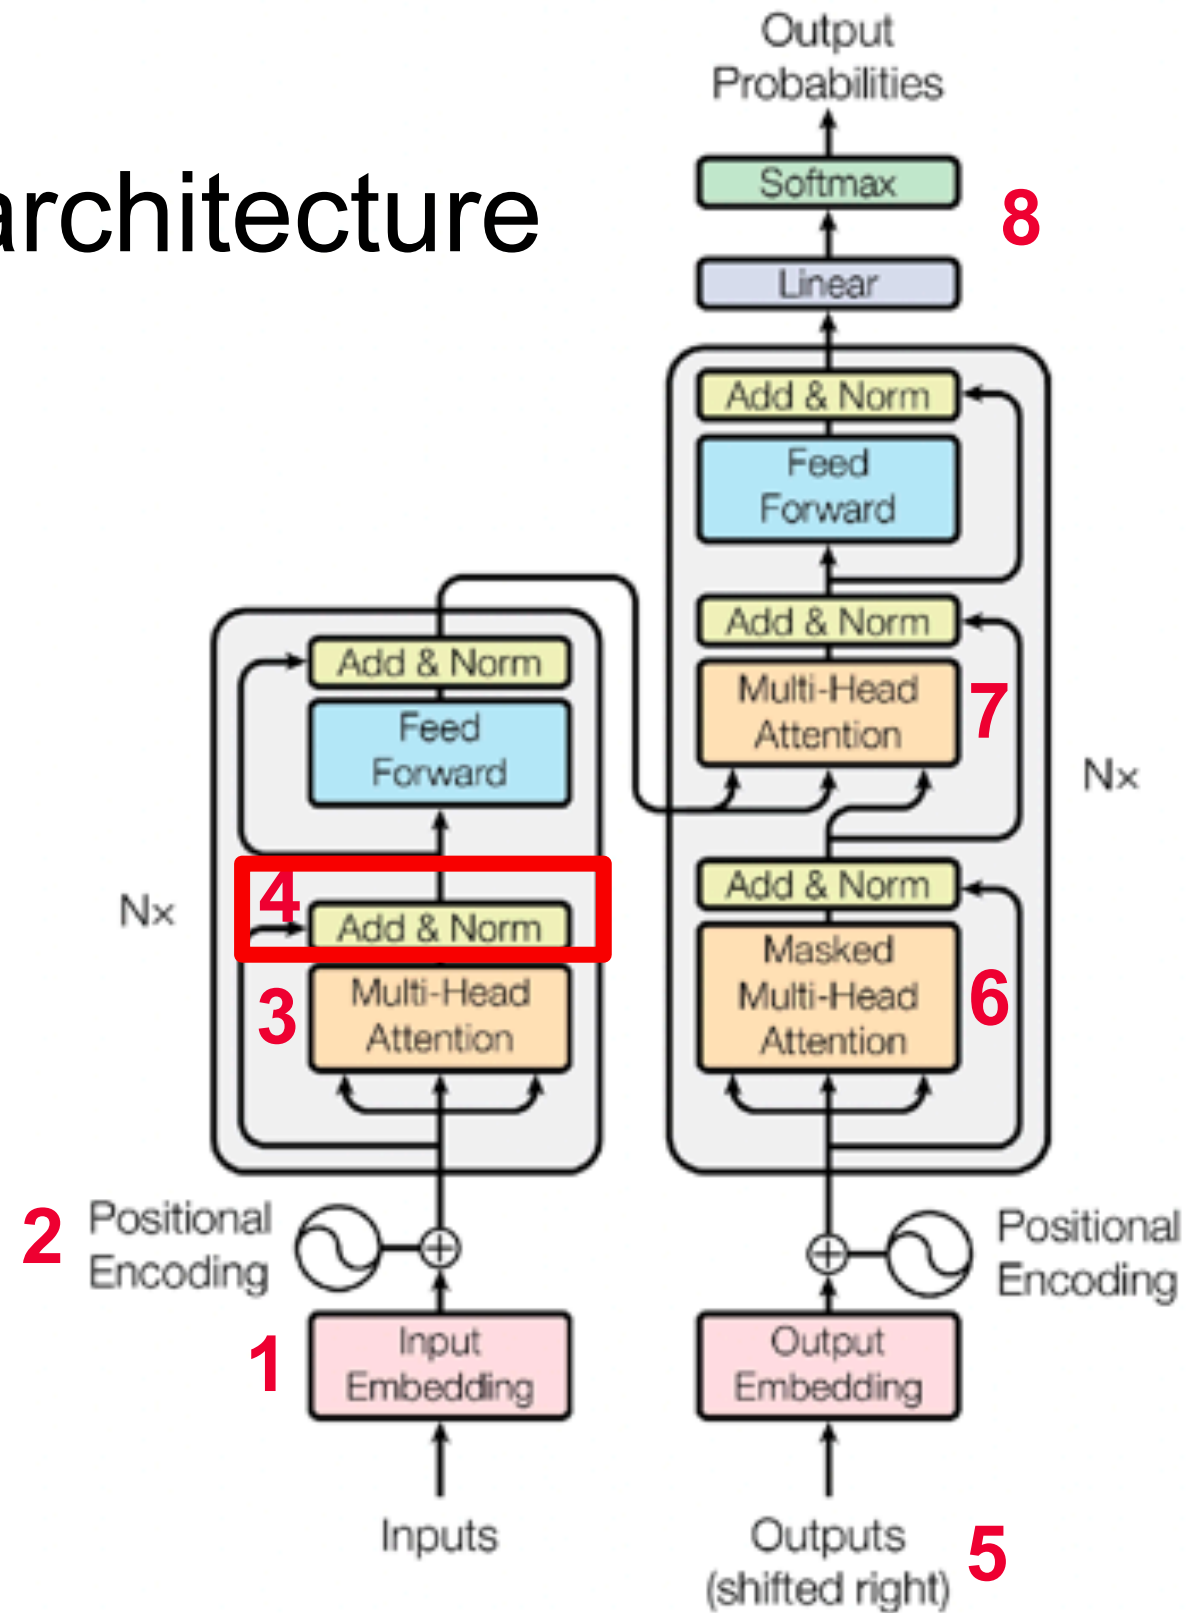




Figure 1: The Transformer - model architecture.


Residual Connections

- Network degradation problem



Layer Normalization

- To get stable gradients and faster convergence
- Normalization(with learnable params) **over the layer(sequence)**
- Work well on **arbitrary** length of sequence



Transformer architecture

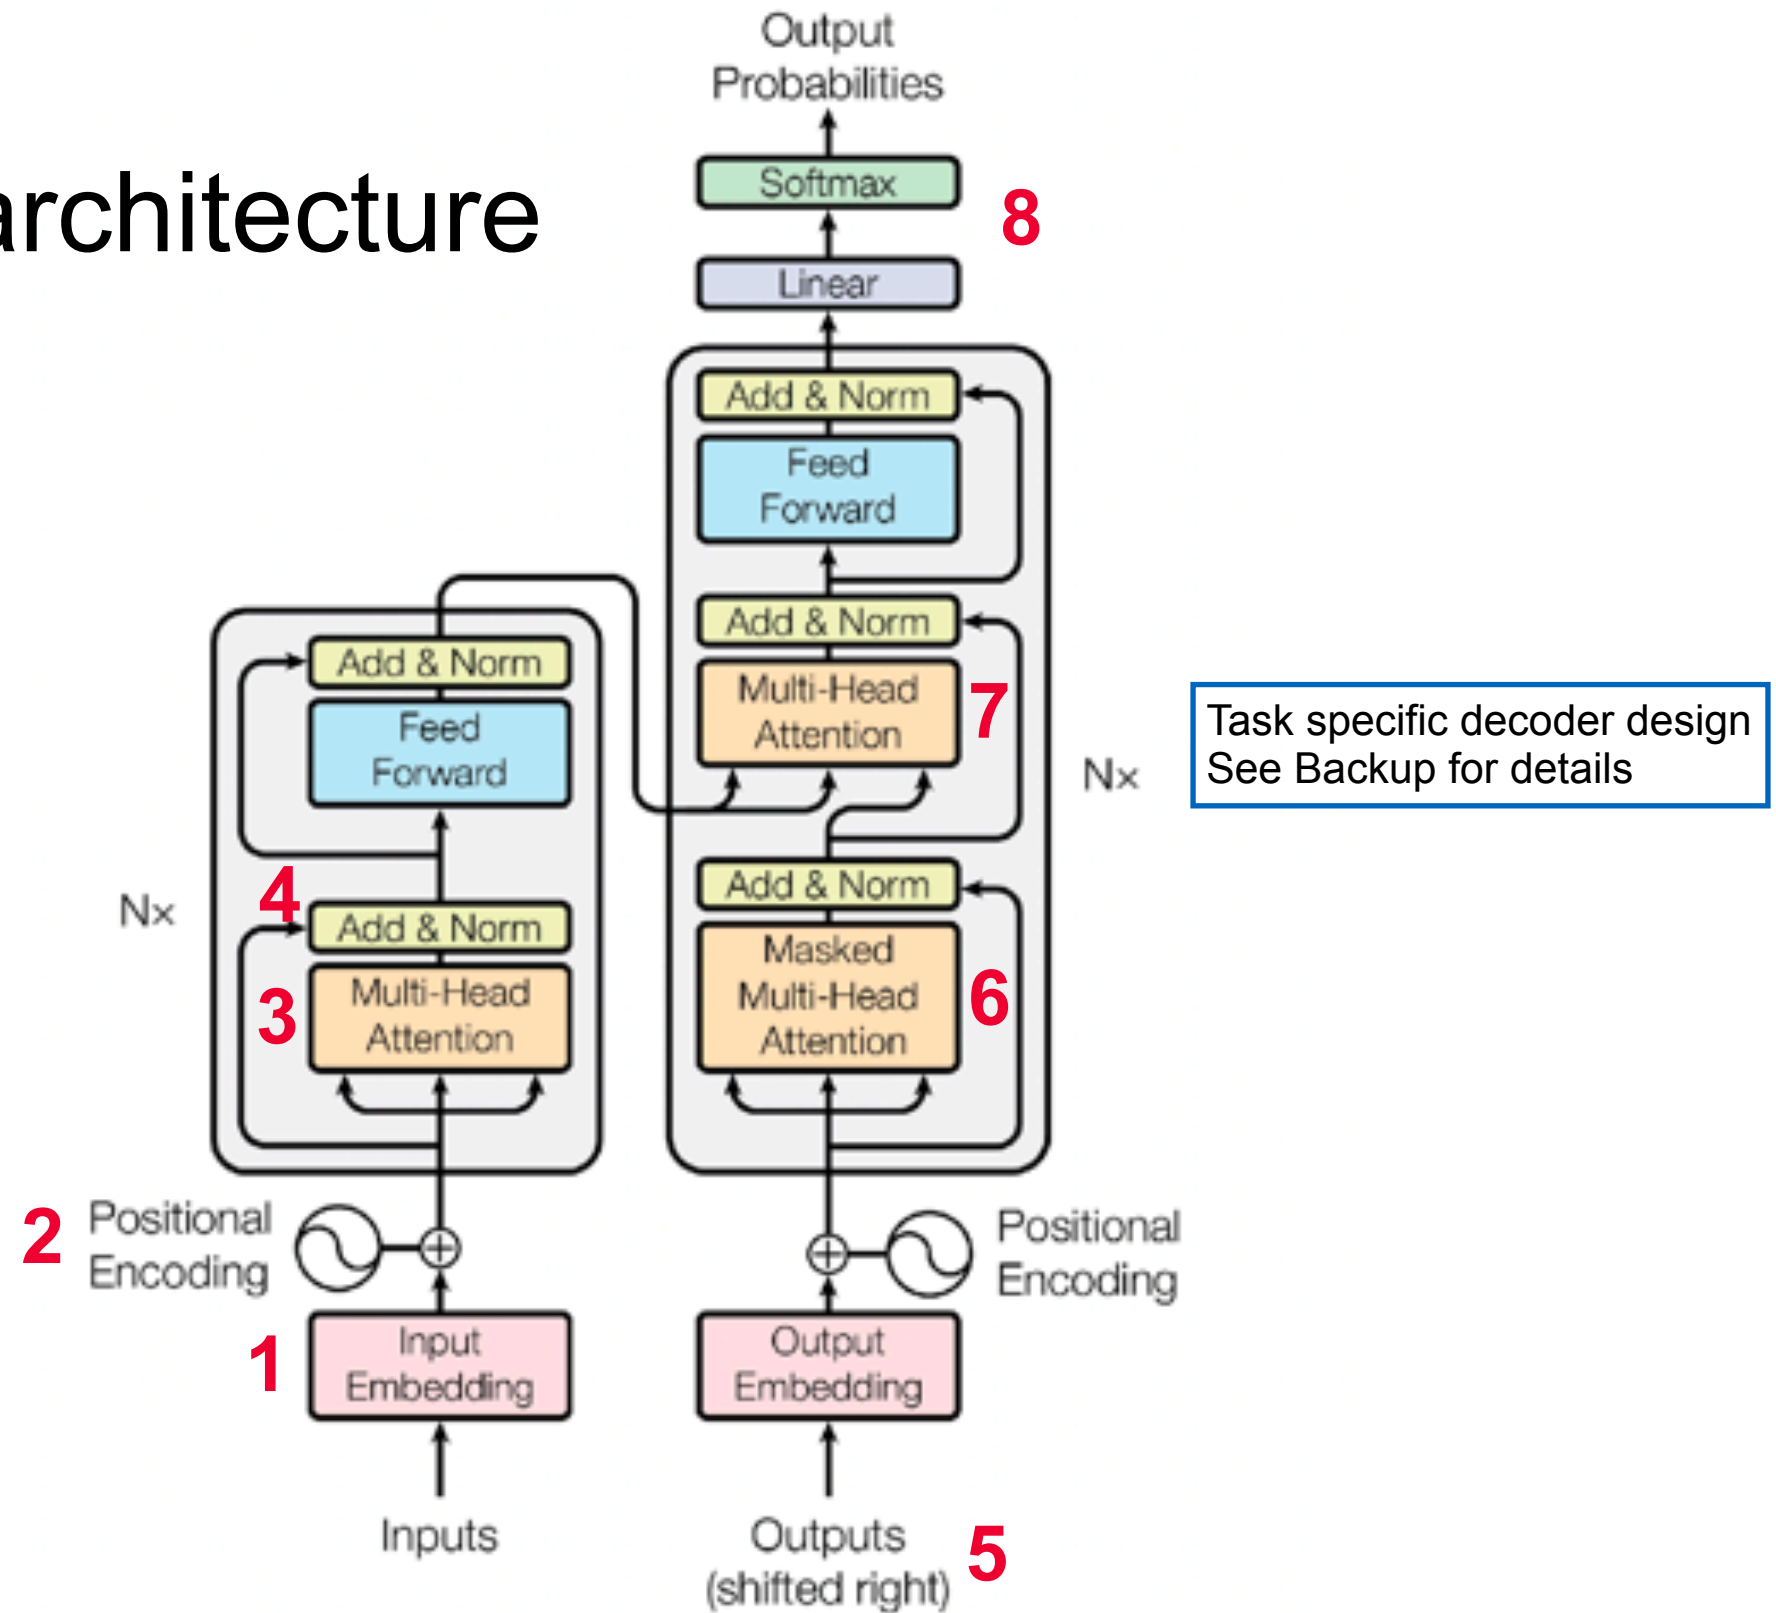


Figure 1: The Transformer - model architecture.

Demo



<https://colab.research.google.com/drive/1hXIQ77A4TYS4y3UthWF-Ci7V7vVUoxmQ?usp=sharing>

VisionTransformer

An Image is Worth 16x16 Words:
Transformers for Image Recognition at Scale

Motivation

- Can we directly apply Transformers for image classification?
- What's the network's focus when classifying images?



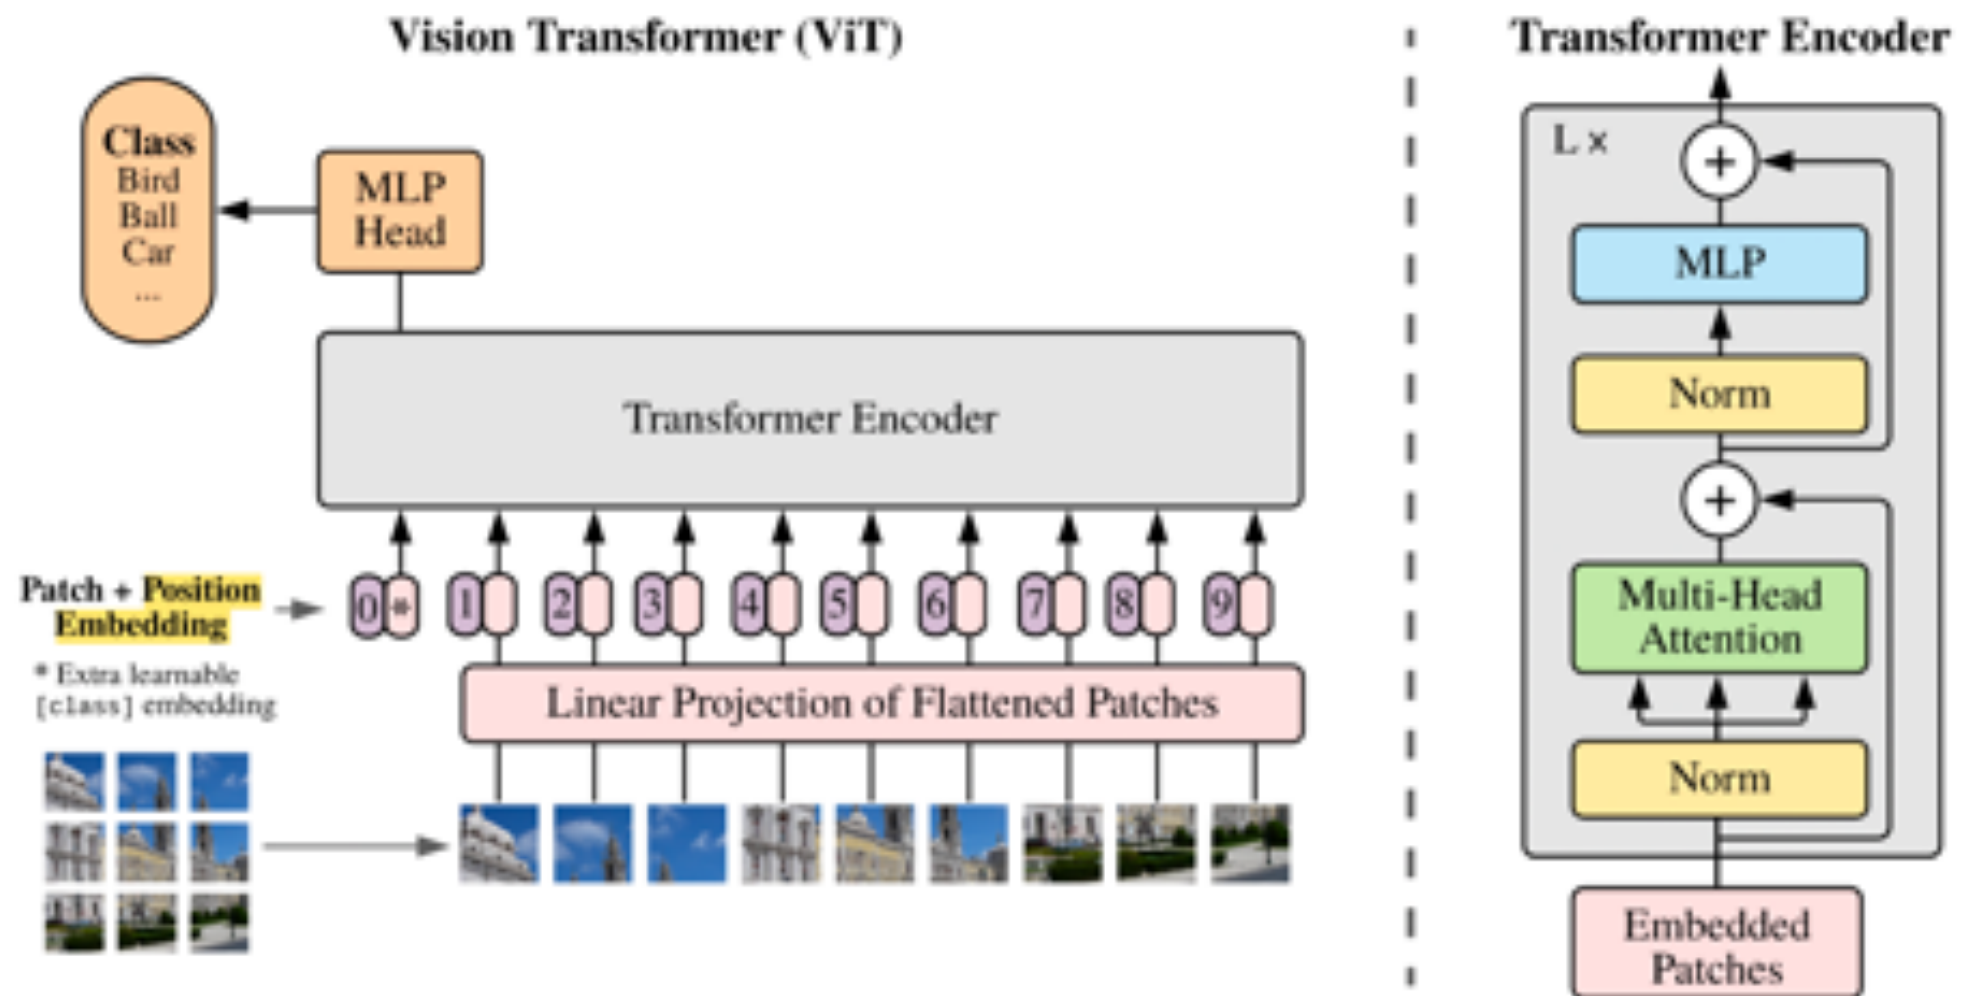
Divide Images into Patches, then Linear Project

- Number of pixels(e.g. 224x224, 600x800...) are much larger than the length of text sequences
- Given an image $x \in \mathbb{R}^{H \times W \times C}$, divide it by 16x16 patches to get N patches: $x_p \in \mathbb{R}^{N \times (P^2 \times C)}$
- Flatten each patch, map to D dimension with trainable linear projection
 - patch embeddings $x_E \in \mathbb{R}^{N \times (1 \times D)}$



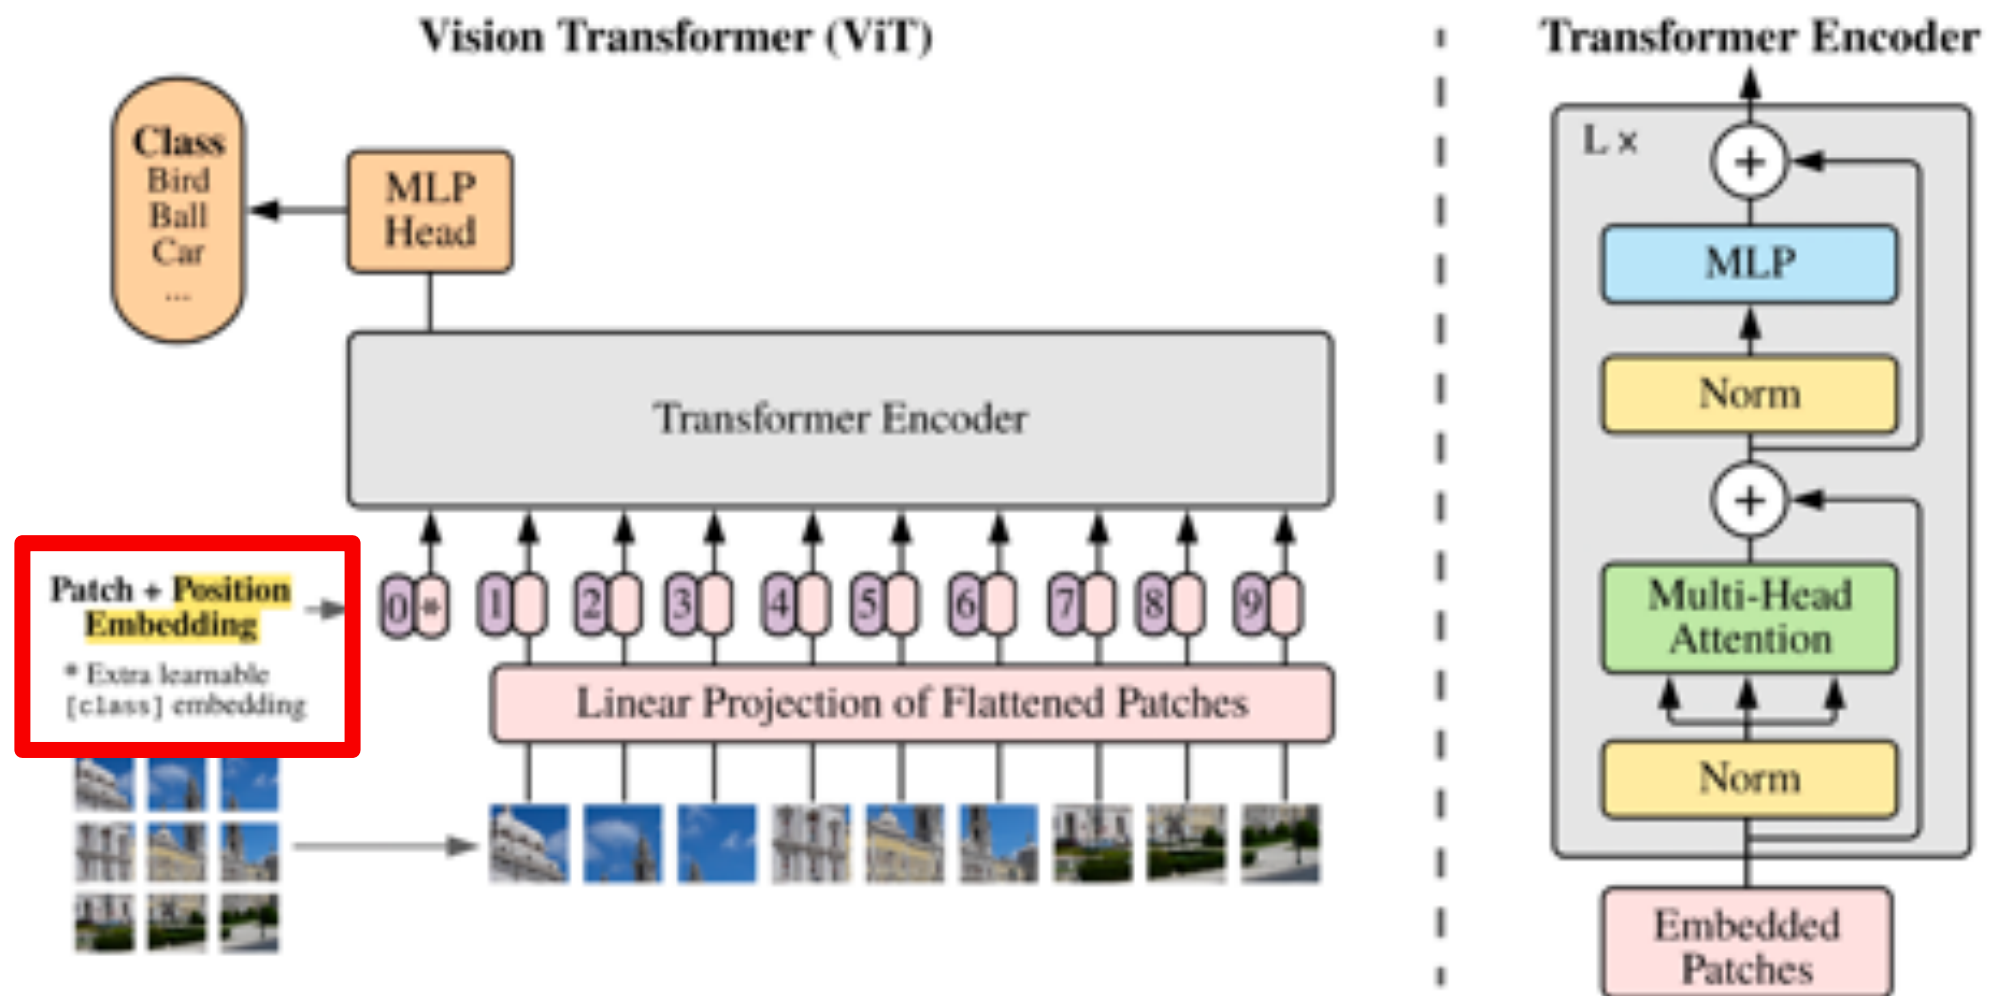
Model Architecture

- Only the Transformer Encoder are used, same design as the original Transformer
- Use patch to divide input images, then project into sequences. Use position embedding
- Task specific MLP head.



Model Architecture

- Only the Transformer Encoder are used, same design as the original Transformer
- Use patch to divide input images, then project into sequences. Use position embedding
- Task specific MLP head.



Position Embeddings (but learnable!)

Different strategies, same goal: to encode the position of each patch!

- **1-dimensional positional embedding(used)**: inputs as a sequence of patches
- **2-dimensional positional embedding**: inputs as a grid of patches in two dimensions
- **Relative positional embeddings**: self-attention between patches for the relative distance

*“In patch-level inputs, the spatial dimensions are **much smaller** than the original pixel-level inputs, e.g., 14×14 as opposed to 224×224 , and learning to represent the spatial relations in this resolution is **equally easy** for these different positional encoding strategies.”*

The [CLS] token

- Self-attention is patch-wise. **BUT how to represent the whole image?**
 - Average all the embeddings
 - Attach an extra learnable [CLS] token, involved in all self-attention calculations




Figure 9: Comparison of class-token and global average pooling classifiers. Both work similarly well, but require different learning-rates.


Demo

<https://colab.research.google.com/drive/1rTz-uwb-nYXAQb6fJq1nWD43ahllwnFU?usp=sharing>

Properties of Vision Transformers

Vision Transformers are data-hungry

- On mid-sized datasets (e.g. ImageNet), poorer performances than ResNets family
- On large datasets (e.g. ImageNet 21K, JFT-300M), state-of-the-art



ResNets perform better with smaller pre-training datasets,
but plateau sooner than ViT for larger datasets

Properties of Vision Transformers


Vision Transformers lack inductive biases in CNNs



Inductive bias: maybe this is a cow


Properties of Vision Transformers

Vision Transformers lack inductive biases in CNNs



Locality

Assume neighbor contents have similar information



Translational equivalence

Translations in input also change the output

Properties of Vision Transformers

Vision Transformers can be easily scaled up

- Stack multiple attention layers
- Use the [CLS] token with task-specific heads

Model	Layers	Hidden size D	MLP size	Heads	Params
ViT-Base	12	768	3072	12	86M
ViT-Large	24	1024	4096	16	307M
ViT-Huge	32	1280	5120	16	632M

Table 1: Details of Vision Transformer model variants.


Applications

What can Transformers do for medical images?


[1] Shamshad, Fahad, et al. "Transformers in medical imaging: A survey." *arXiv preprint arXiv:2201.09873* (2022).

[2] <https://github.com/fahadshamshad/awesome-transformers-in-medical-imaging>

Multi-organ Segmentation




(a) From left to right: Input, ResUNet (CNN), CBAM (CNN), Dual Attention (CNN), and UTNet (ViT). The outline indicates the ground-truth annotation.




(b) From left to right: Input, Ground truth, UNETR (ViT), CoTr (ViT), TransUNet (ViT), and nnUNet (CNN). Numbers are Dice scores.

Vessel Segmentation




(c) First row: Ground truth. Second row (left) CNN based approach. Second row (right) ViT based approach.

Low-Dose Image Enhancement



d) Top to bottom and left to right: Normal Dose CT, Low Dose CT, Non-Local Mean (hand-crafted), RED-CNN (CNN), MAP-NN (CNN), and TransCT (ViT)

Image Reconstruction



(e) From left to right (top row): Fourier method, GAN (CNN), SAGAN (CNN), and SLATTER (ViT). Bottom row shows corresponding error maps.

Thanks for your ATTENTION :)

Questions?

Backup

Batch Normalization

Input: Values of x over a mini-batch: $\mathcal{B} = \{x_1 \dots x_m\}$;
 Parameters to be learned: γ, β
Output: $\{y_i = \text{BN}_{\gamma, \beta}(x_i)\}$




$$\mu_{\mathcal{B}} \leftarrow \frac{1}{m} \sum_{i=1}^m x_i \quad // \text{ mini-batch mean}$$

$$\sigma_{\mathcal{B}}^2 \leftarrow \frac{1}{m} \sum_{i=1}^m (x_i - \mu_{\mathcal{B}})^2 \quad // \text{ mini-batch variance}$$

$$\hat{x}_i \leftarrow \frac{x_i - \mu_{\mathcal{B}}}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}} \quad // \text{ normalize}$$


$$y_i \leftarrow \gamma \hat{x}_i + \beta \equiv \text{BN}_{\gamma, \beta}(x_i) \quad // \text{ scale and shift}$$

- To get stable gradients and faster convergence
- Normalization(with learnable params) **over the batch**
- **Problem:**
 - Sequence lengths vary a lot in mini-batches, causing instability
 - Fail to deal with **unseen** length of input sequence (e.g. testing seq. much longer than training seq.)



Layer Normalization

- To get stable gradients and faster convergence
- Normalization(with learnable params) **over the layer(sequence)**
- Work well on **arbitrary** length of sequence



Transformer architecture

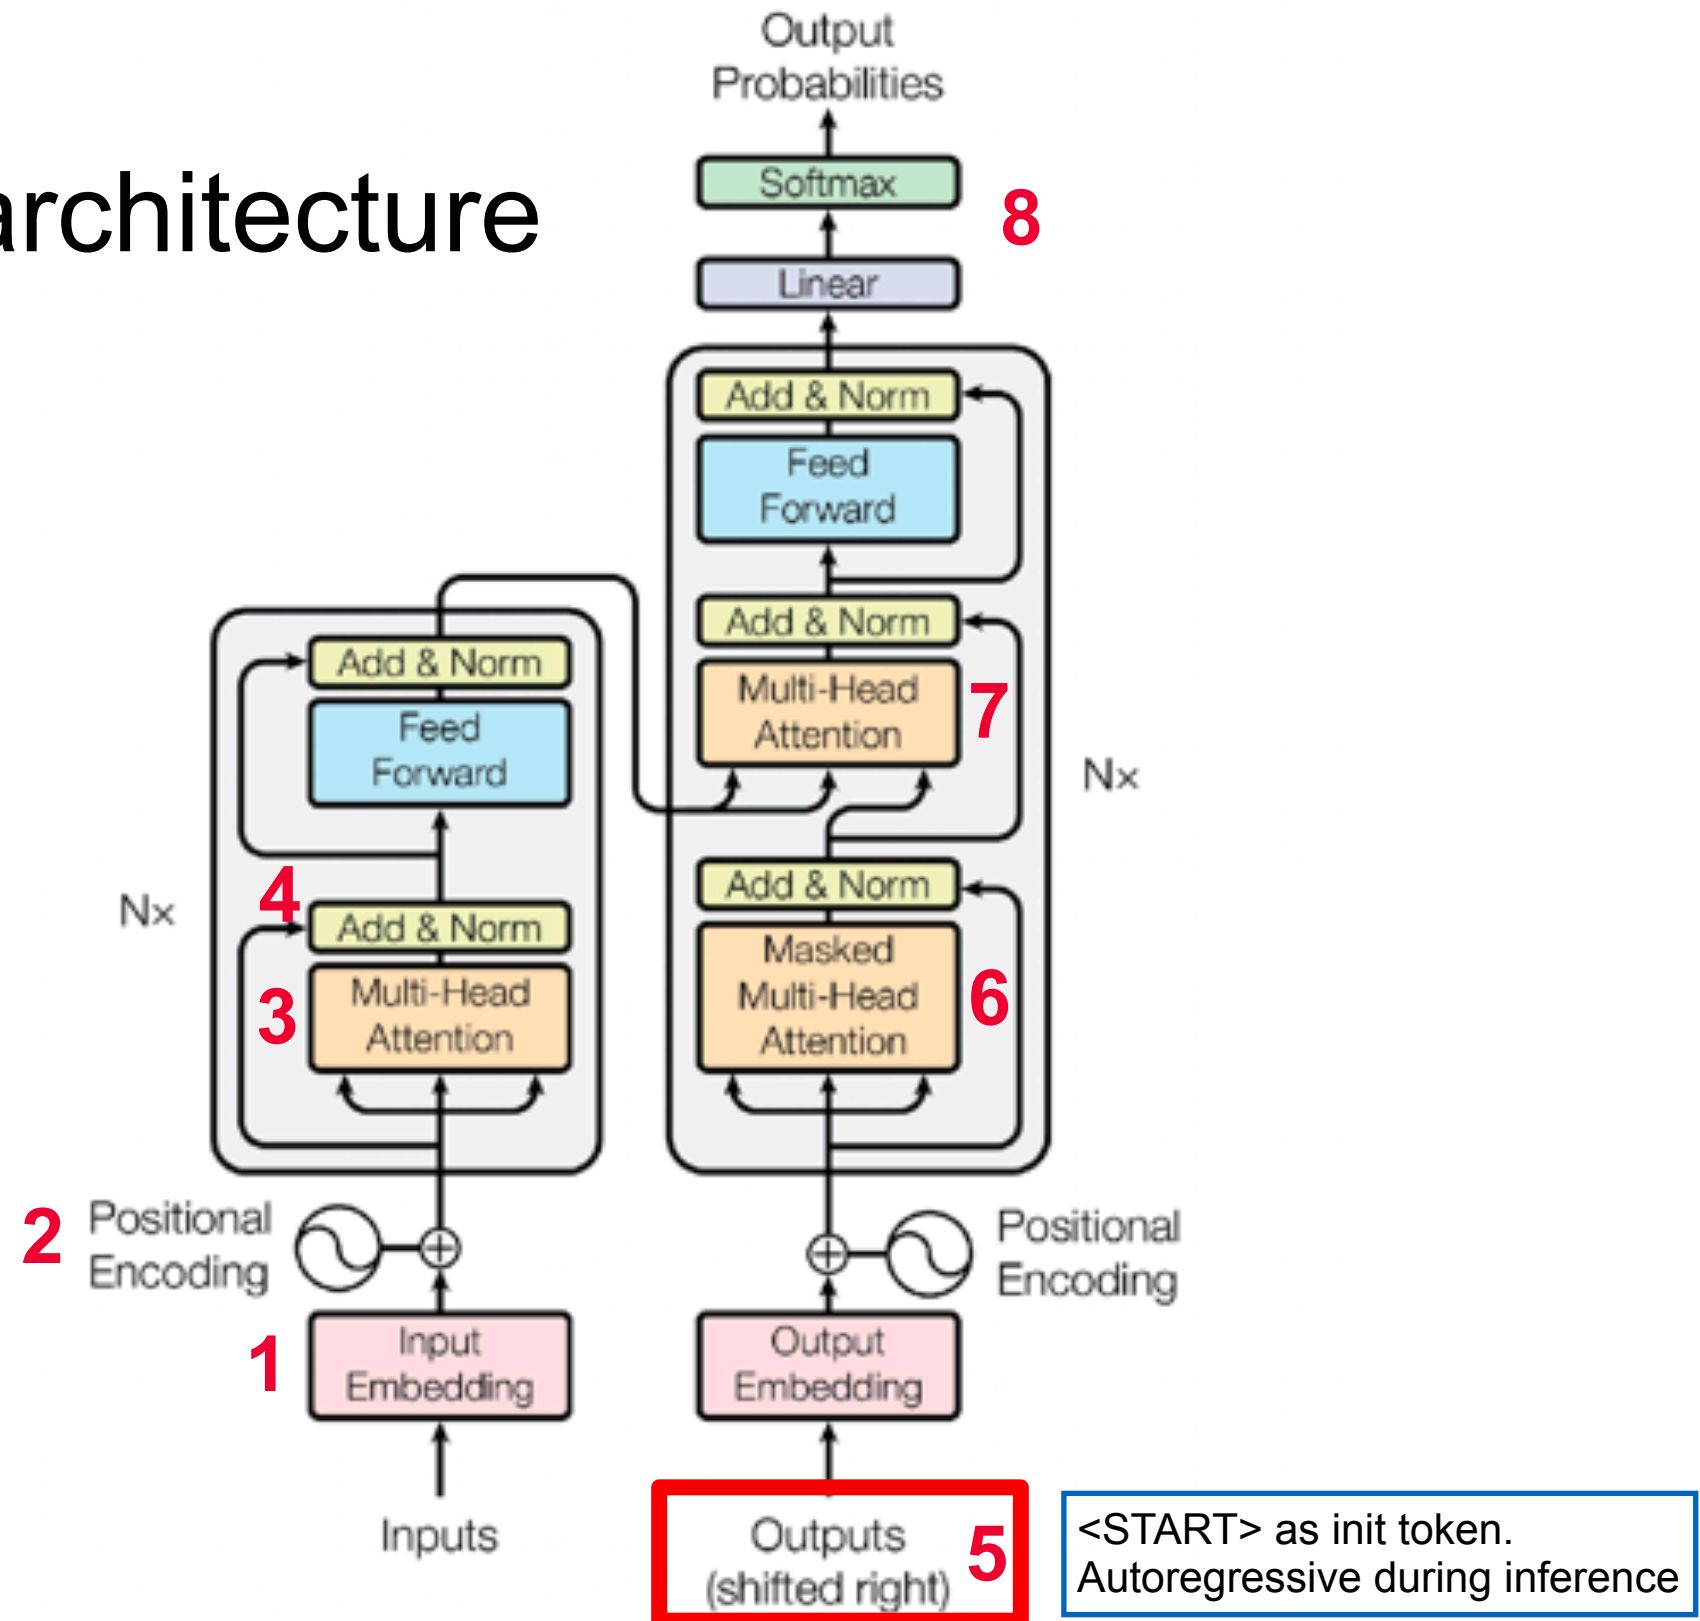
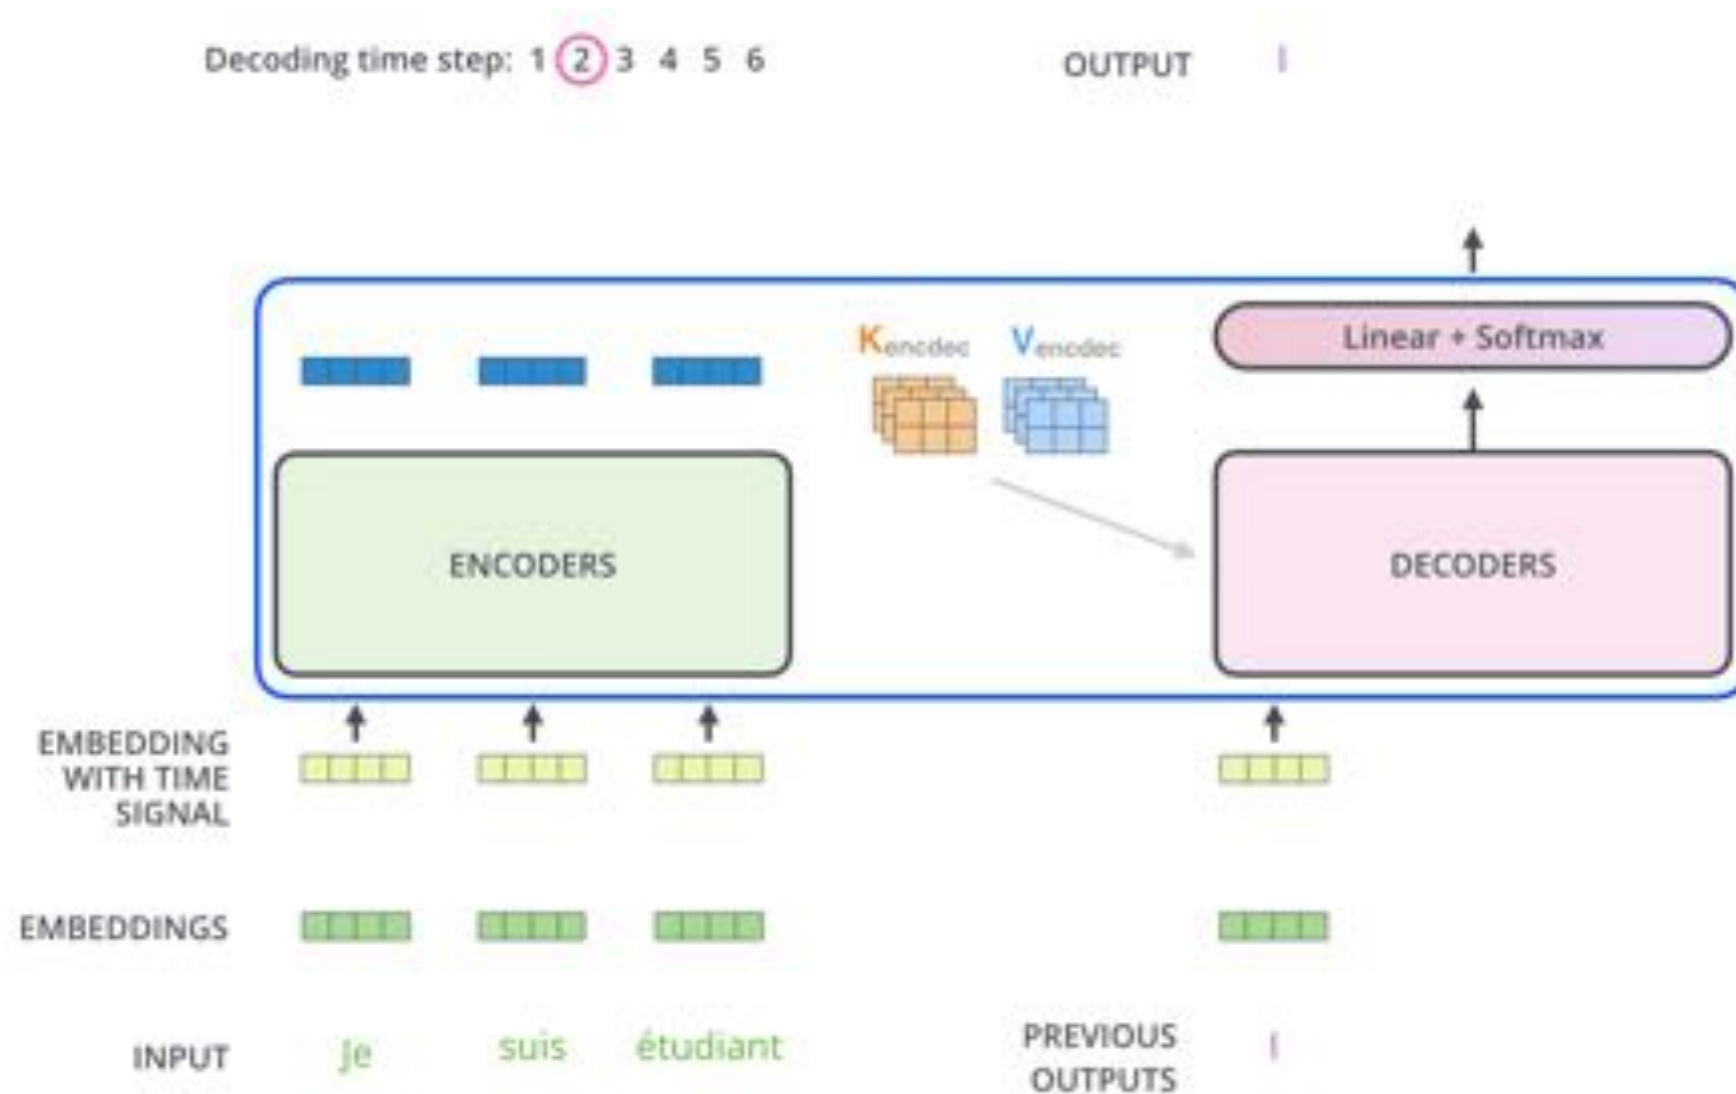


Figure 1: The Transformer - model architecture.

Auto-regressive models

- Auto-regressive: use previous output as current input



Transformer architecture

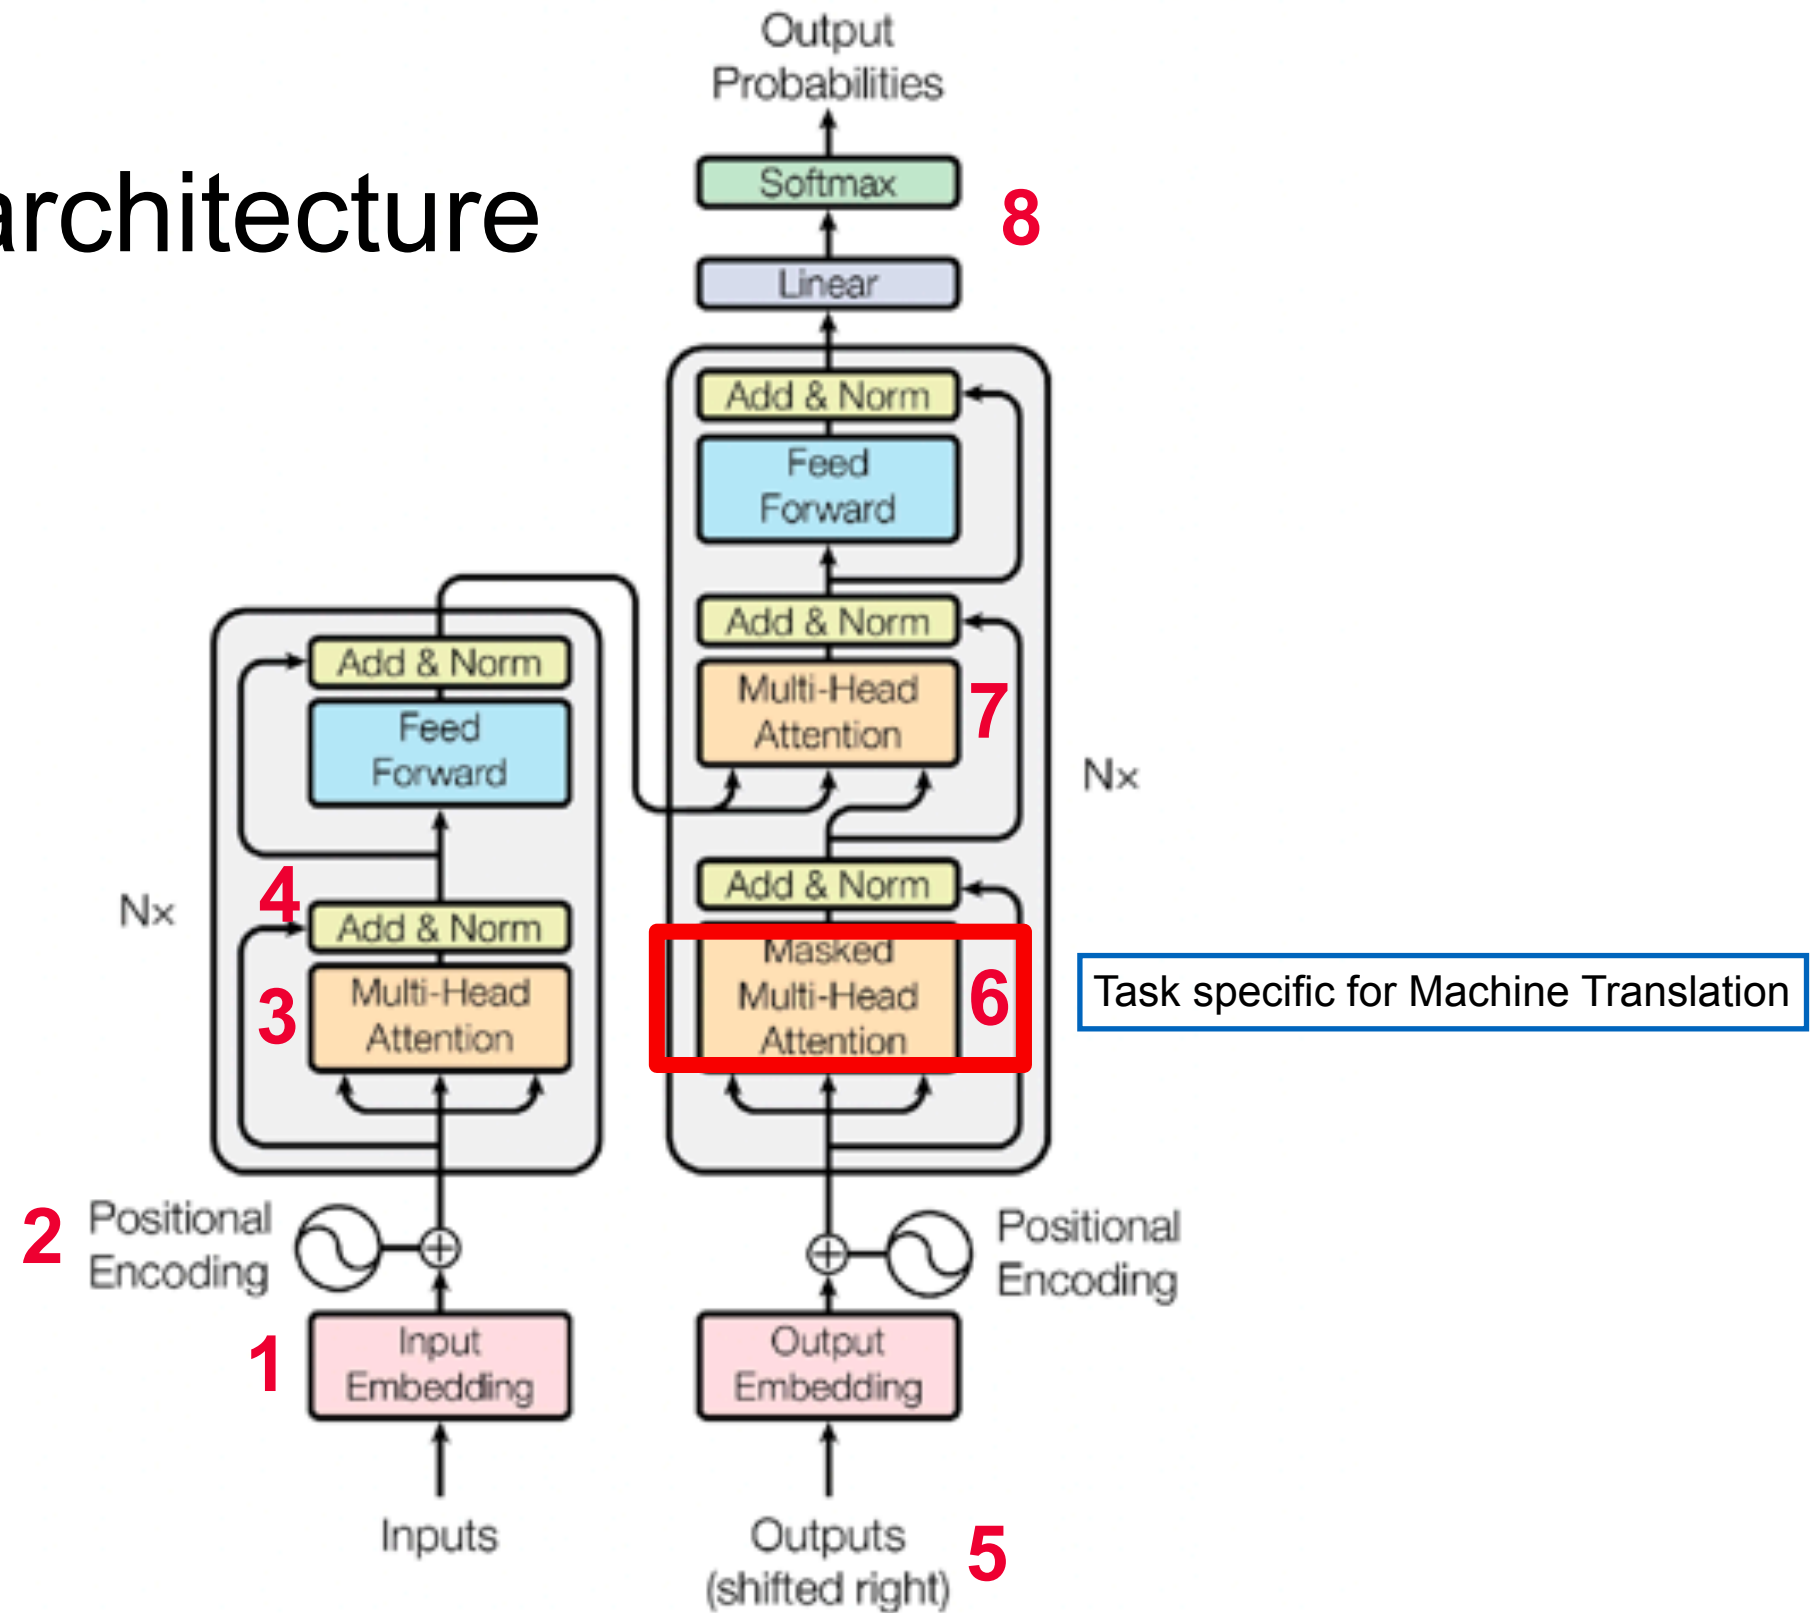


Figure 1: The Transformer - model architecture.

Transformer architecture

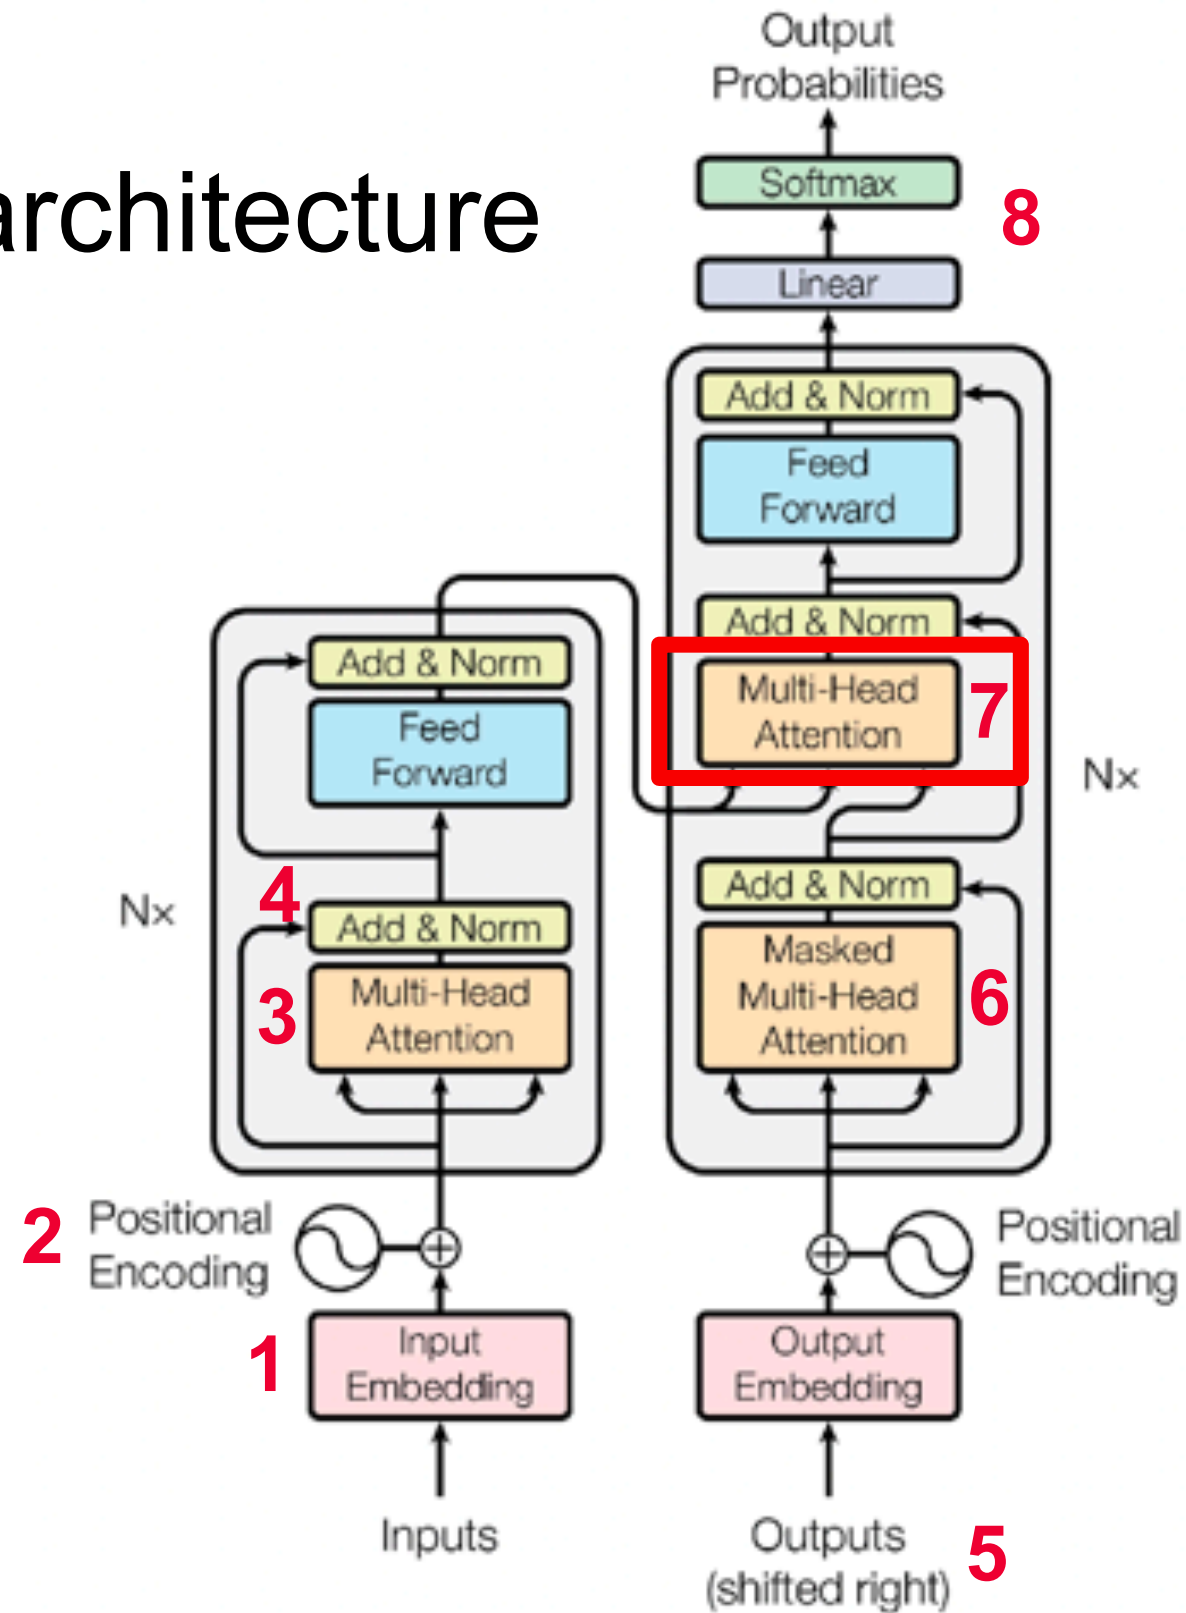
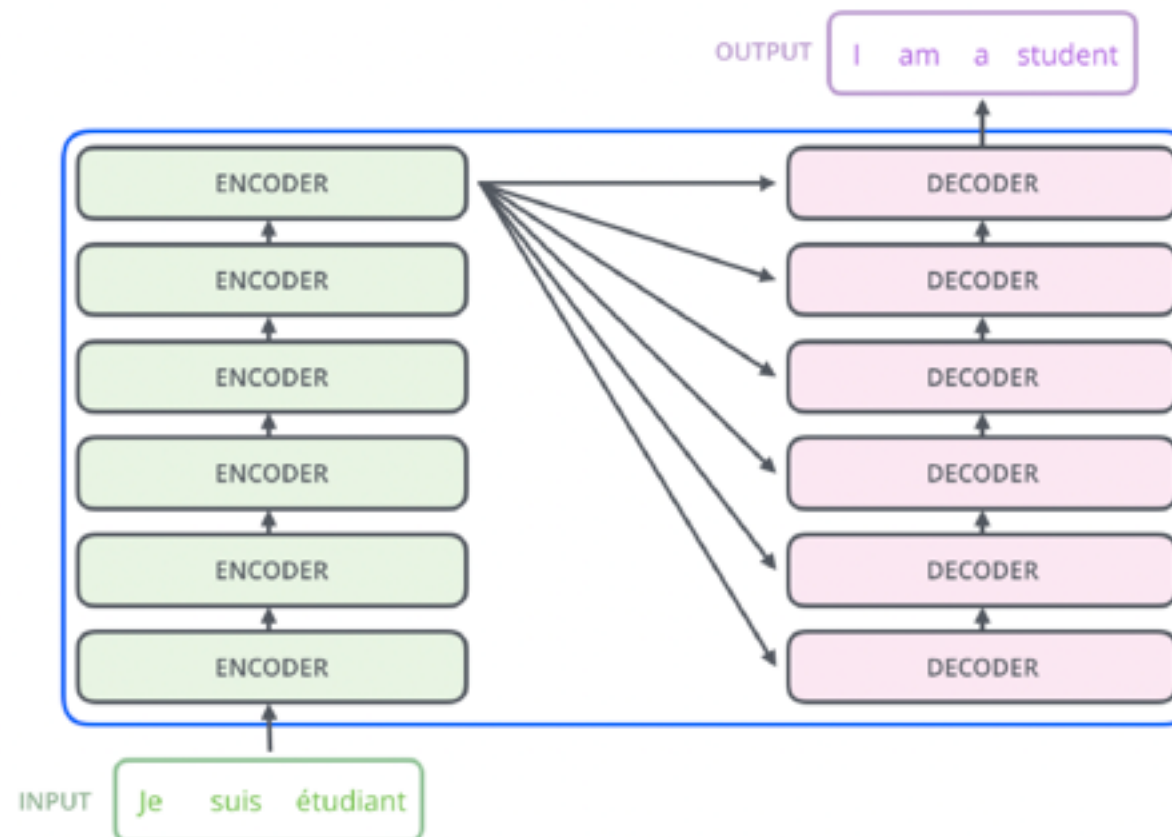


Figure 1: The Transformer - model architecture.

Attention again, but for encoders

- **Keys** and **Values**: from the the last encoder, shared
- **Queries**: from the previous decoder layer



“This allows every position in the decoder to attend over all positions in the input sequence. It mimics the typical encoder-decoder attention mechanisms in sequence-to-sequence models”

Transformer architecture

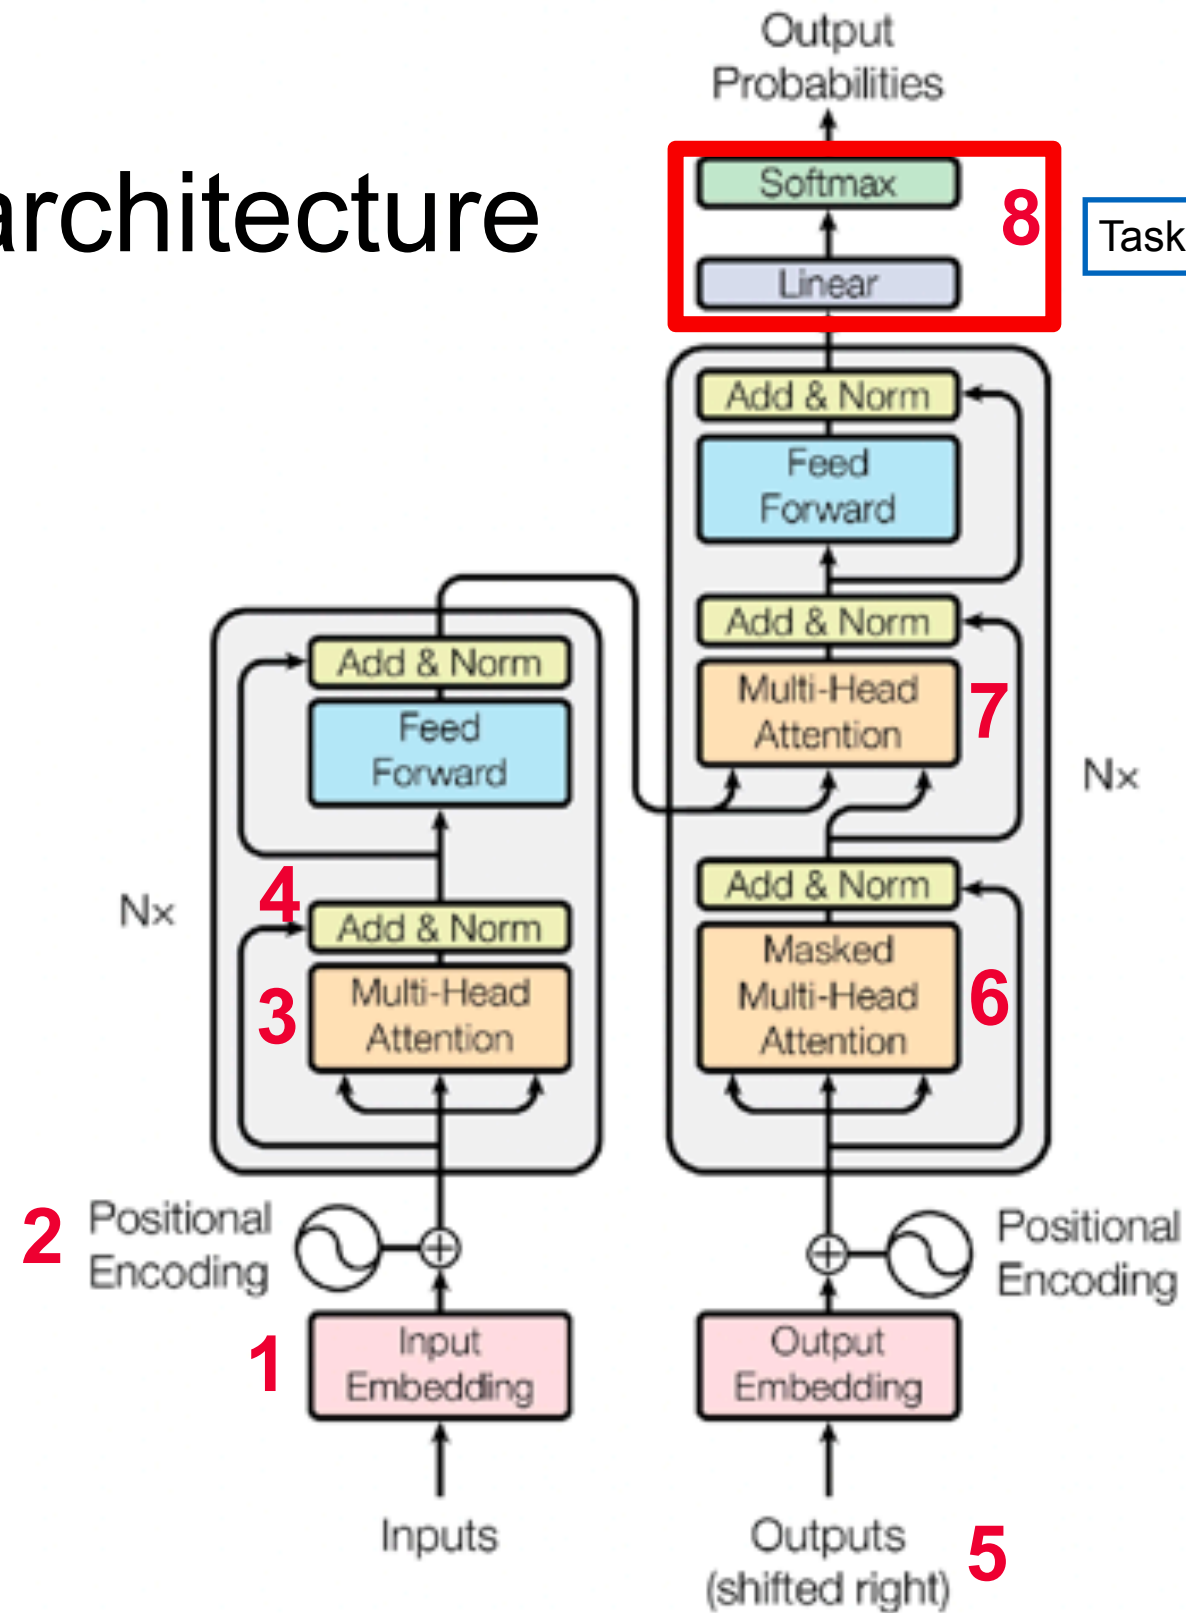



Figure 1: The Transformer - model architecture.

Loss function



- Cross-entropy
- KL Divergence



VisionTransformer

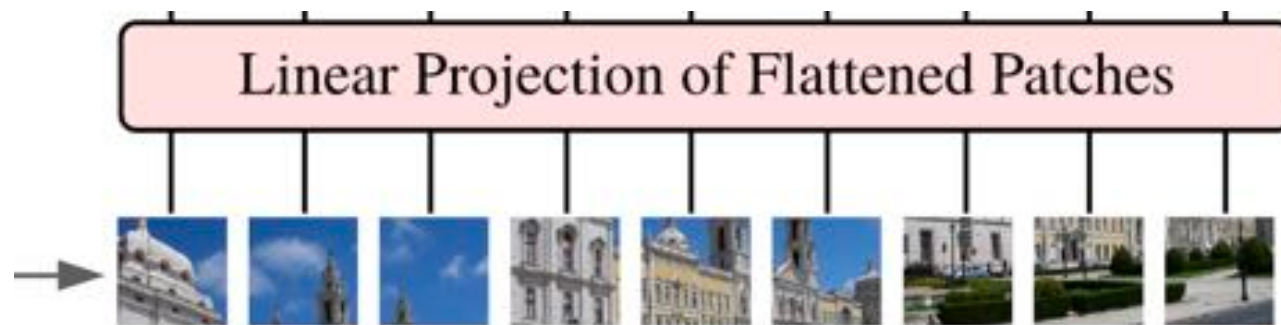
Position Embeddings (but learnable!)

- Hyperparameters can affect learned position embeddings



Divide Images into Patches, then Linear Project

- Given an image $x \in \mathbb{R}^{H \times W \times C}$, divide it by 16x16 patches to get N patches: $x_p \in \mathbb{R}^{N \times (P^2 \times C)}$
- Flatten each patch, map to D dimension with trainable linear projection
 - patch embeddings $x_E \in \mathbb{R}^{N \times (1 \times D)}$
- For a 224x224 image, only 14x14 patches (embeddings)!
- Reduce number of sequence length



Pre-training & Fine-tuning

- Typically, ViTs are pre-trained on large datasets, and fine-tuned to (smaller) downstream tasks.
- Initializing a new task-specific MLP head.
- **Input images have higher resolution than training images (224x224)?**
- Keep patch size same as 16x16
- Larger result sequences than 14x14

Interpolation needed for the position embeddings!

Properties of Vision Transformers

Vision Transformers has less inductive biases than CNNs

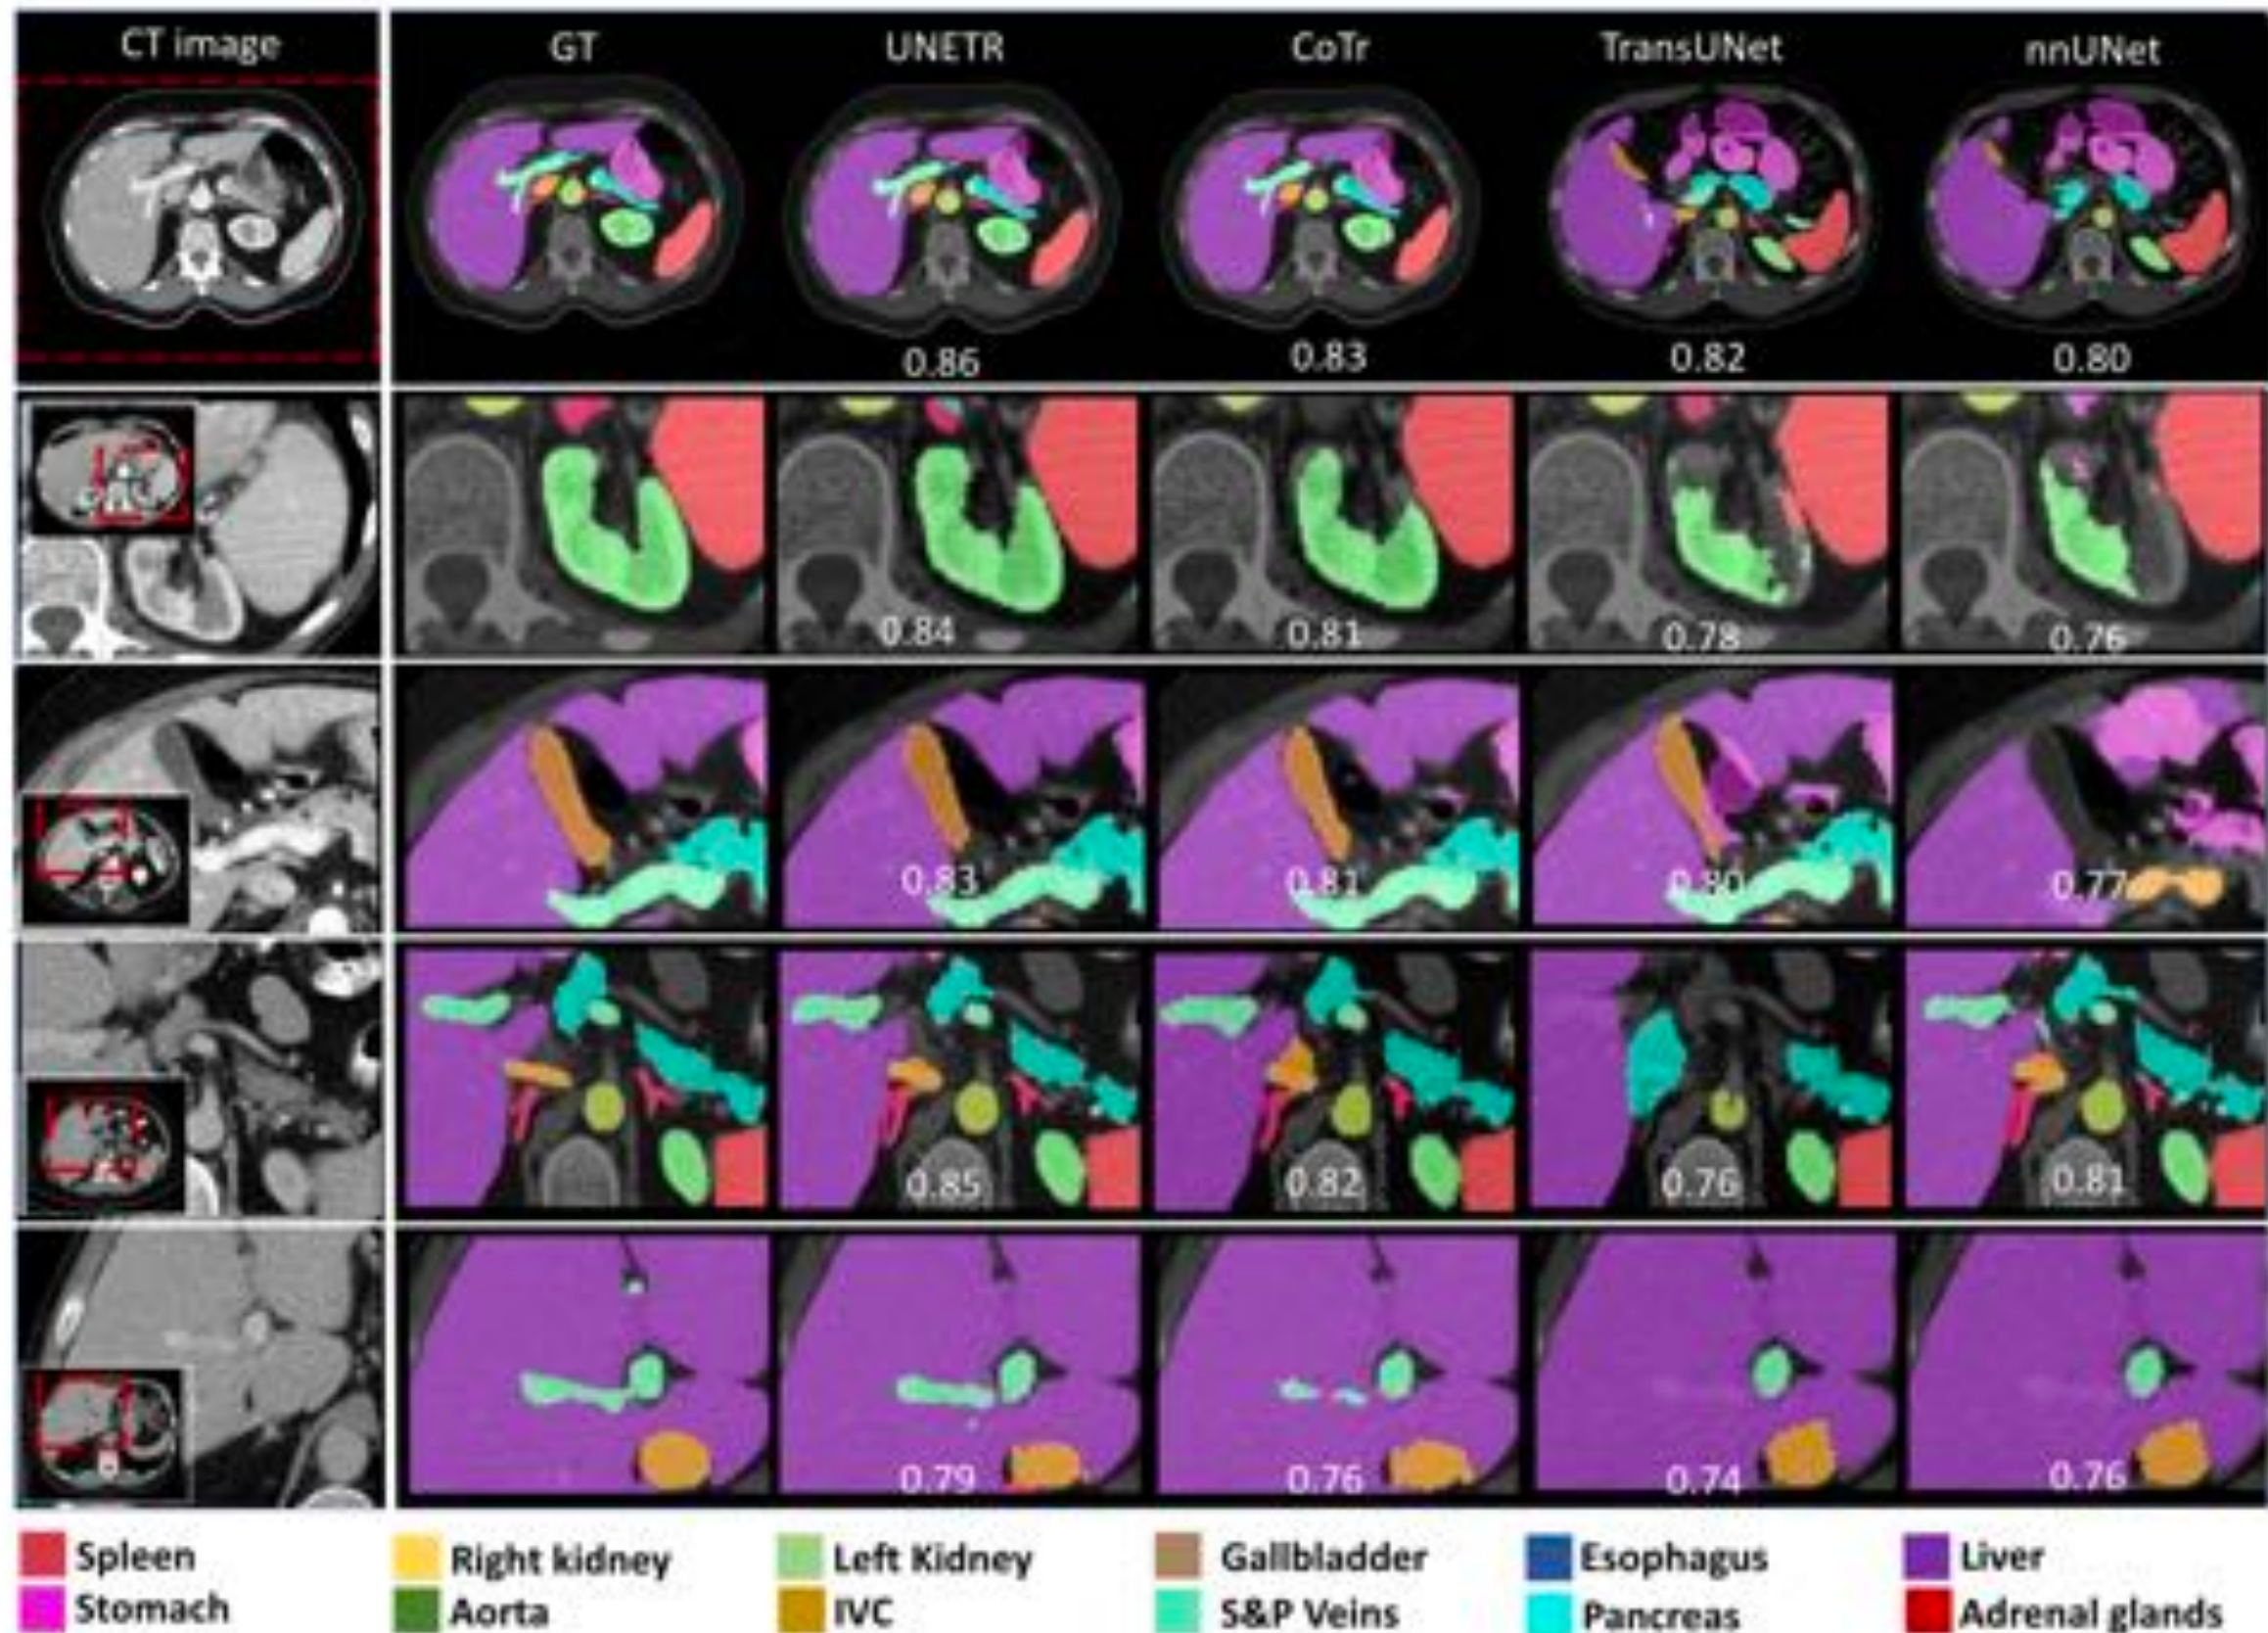
Without inductive biases, Vision Transformers have to learn everything from scratch

- Less data: worse performances than CNNs
- Big data: outperform CNNs

“Transformers lack some of the inductive biases inherent to CNNs, such as translation equivariance and locality, and therefore do not generalize well when trained on insufficient amounts of data... However, we find that large scale training trumps inductive bias.”

Applications

3D Medical Image Segmentation



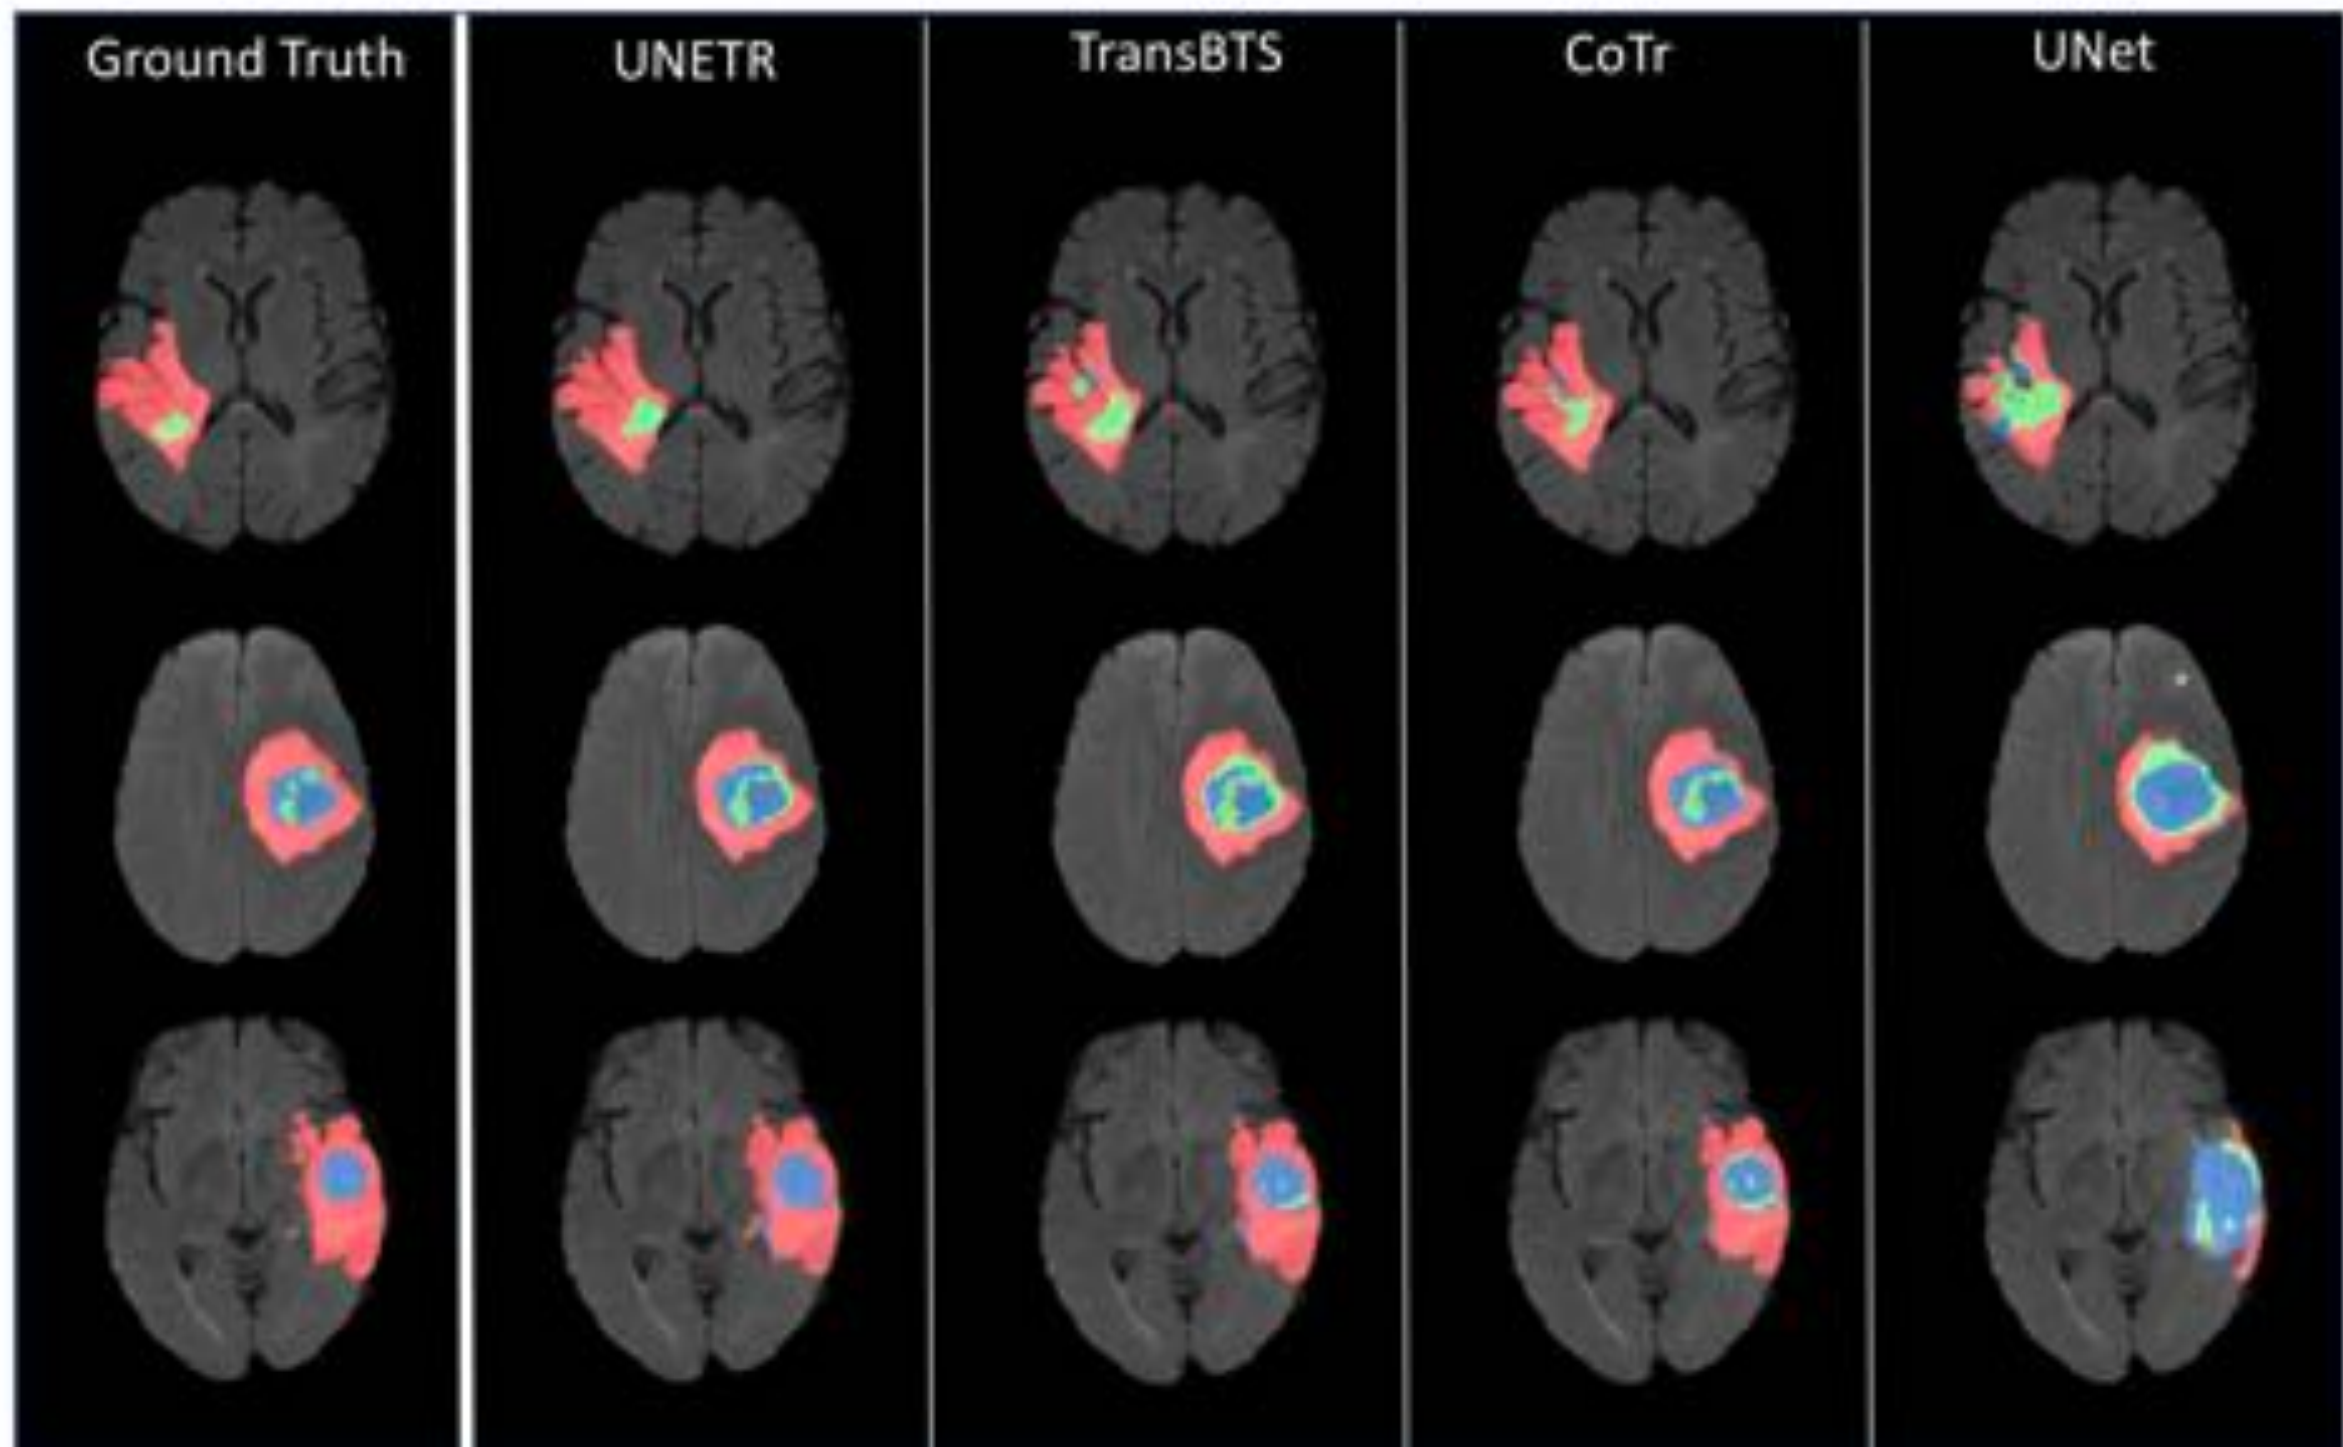


Figure 4. UNETR effectively captures the fine-grained details in segmentation outputs. The Whole Tumor (WT) encompasses a union of red, blue and green regions. The Tumor Core (TC) includes the union of red and blue regions. The Enhancing Tumor core (ET) denotes the green regions.

Model Architecture

- Input 3D medical image data
- VisionTransformer for encoder
- CNN for decoder

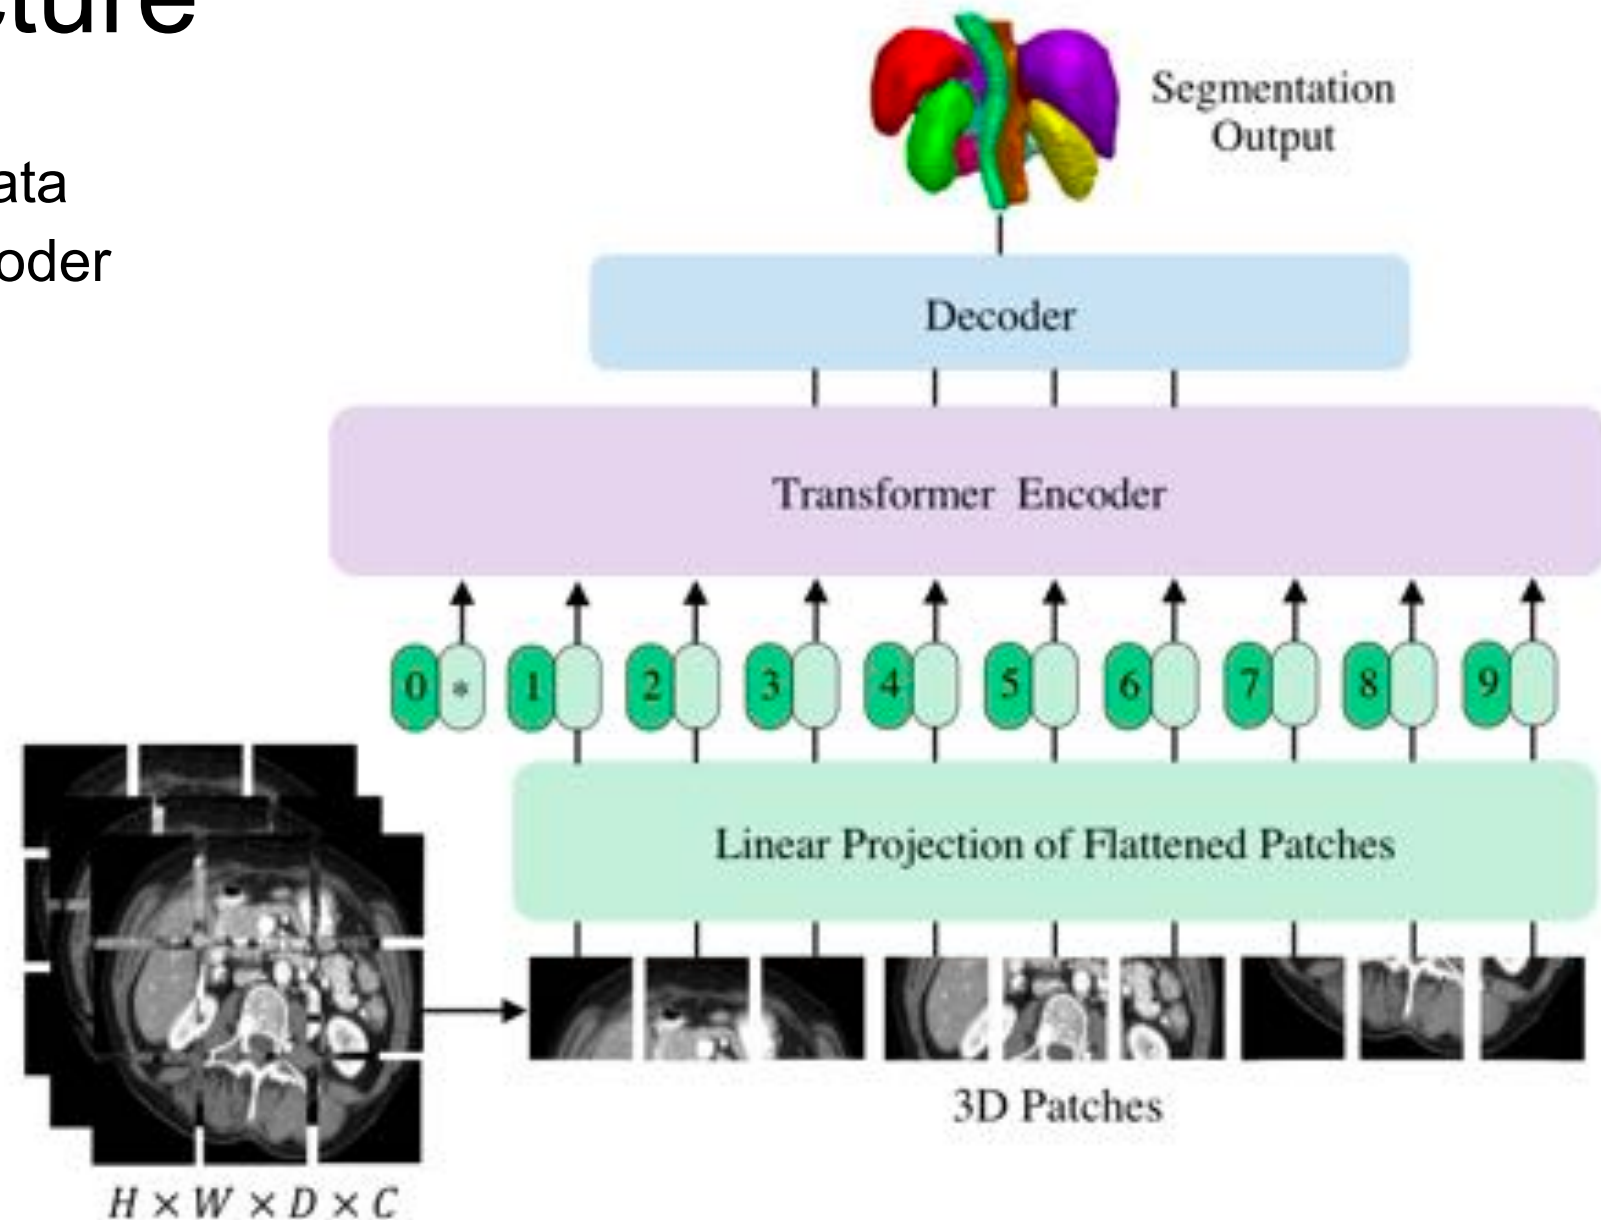
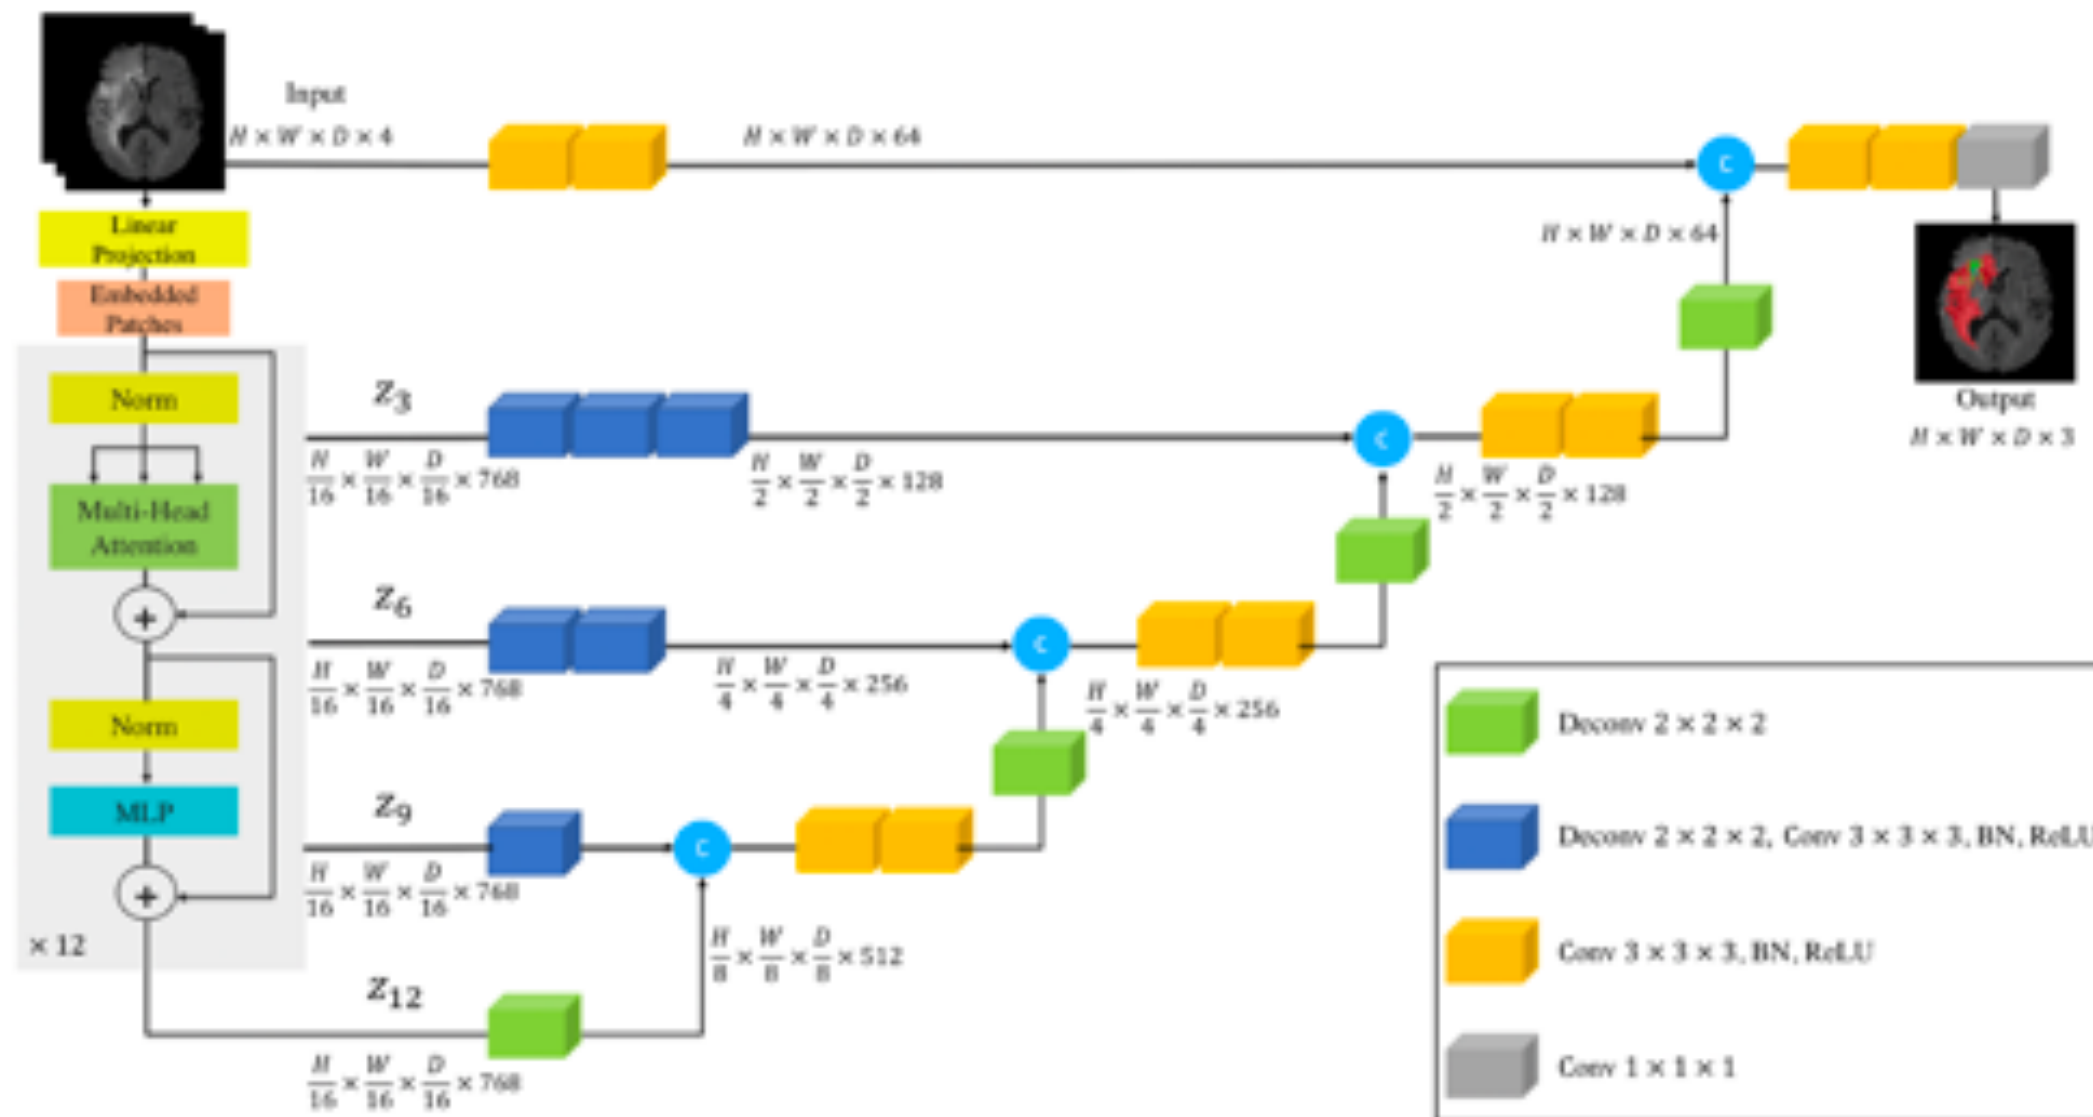


Figure 1. Overview of UNETR. Our proposed model consists of a transformer encoder that directly utilizes 3D patches and is connected to a CNN-based decoder via skip connection.

Model Architecture

- Similar to UNet, use skip connections at the {3,6,9,12}th attention layer
- CNNs as decoder for segmentation




Demo

[https://github.com/Project-MONAI/tutorials/blob/main/3d_segmentation/
unetr_btcv_segmentation_3d.ipynb](https://github.com/Project-MONAI/tutorials/blob/main/3d_segmentation/unetr_btcv_segmentation_3d.ipynb)

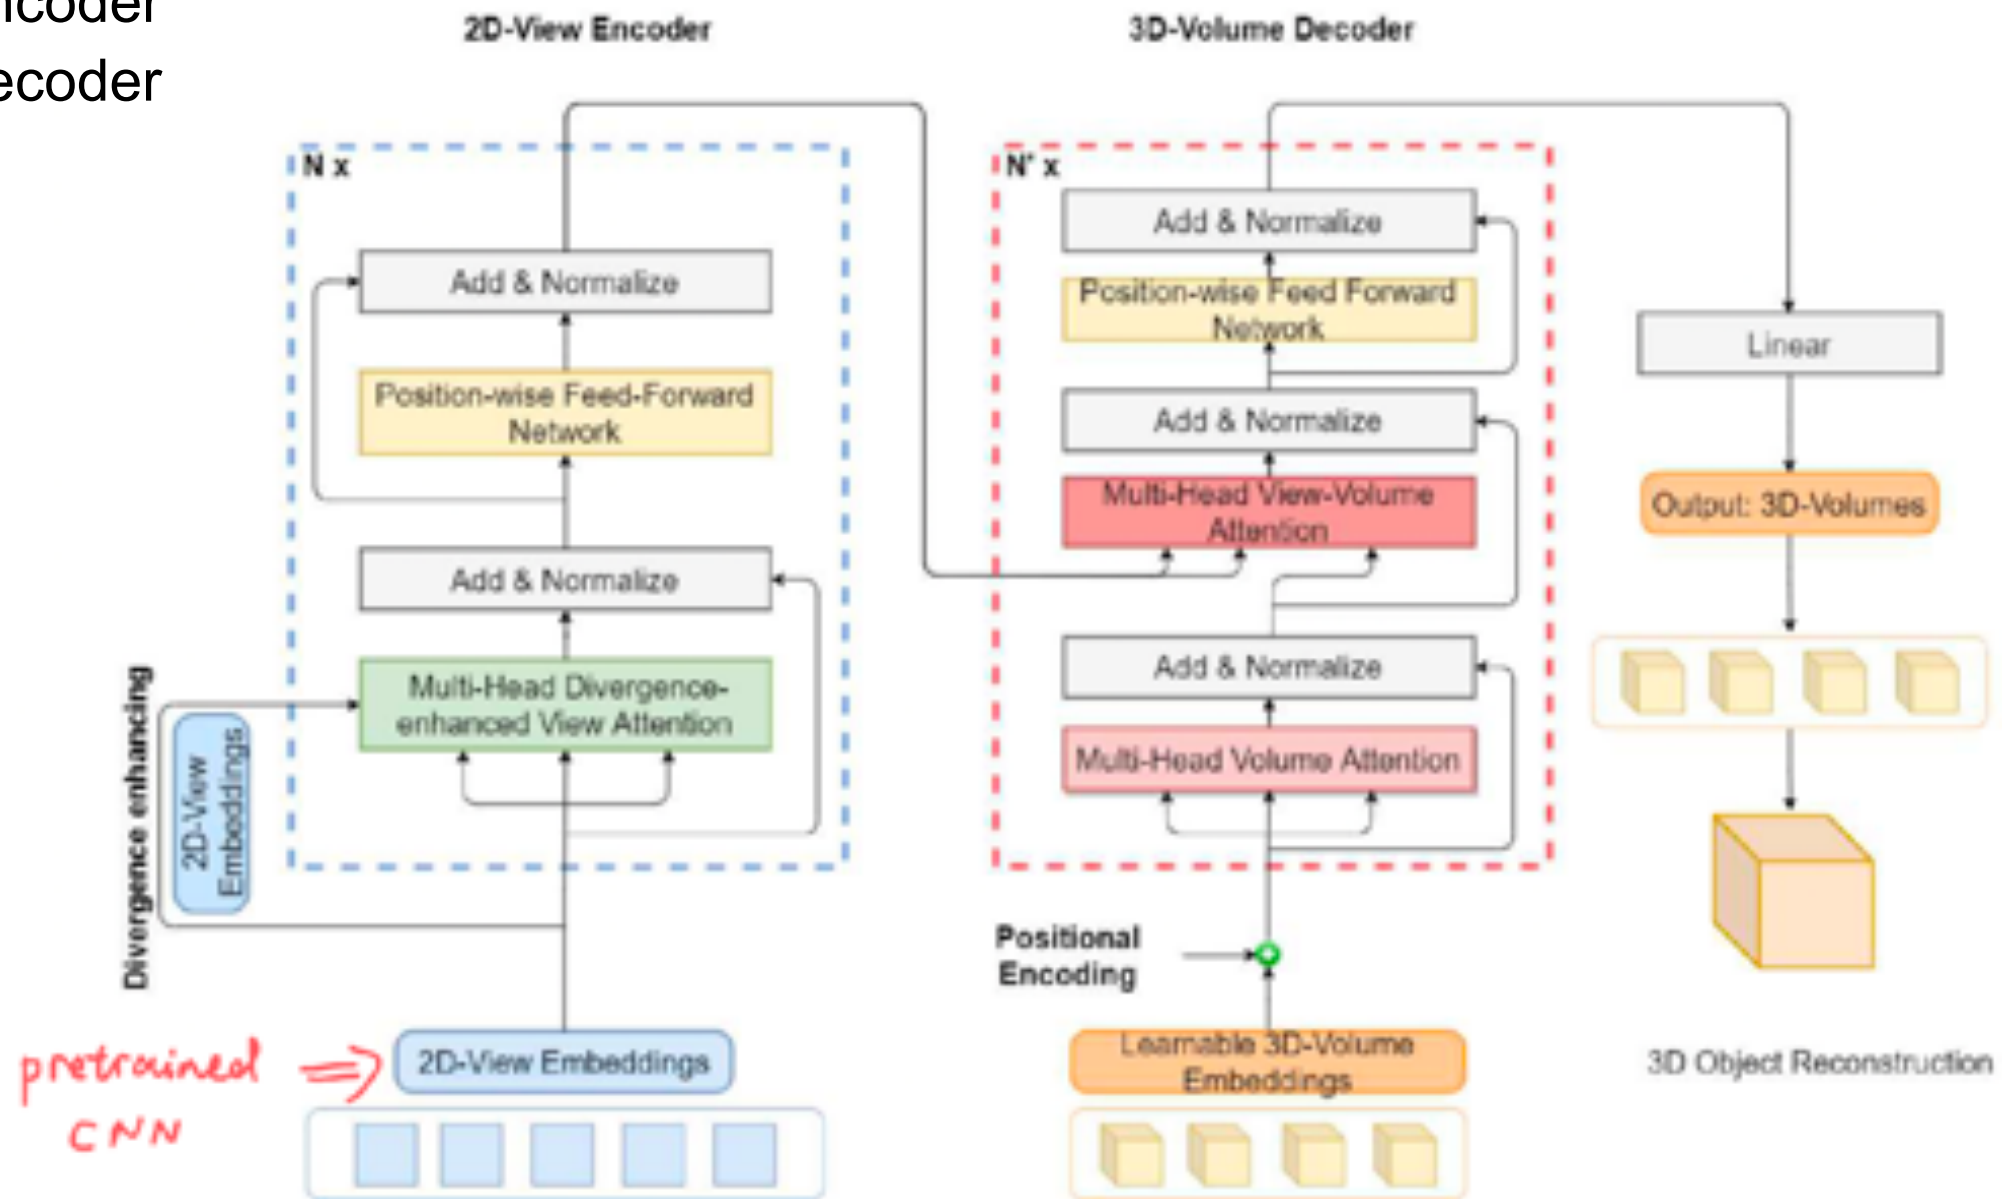
Applications

Multi-view 3D Reconstruction with Transformer



Model Architecture

- 2D ViT encoder
- 3D ViT decoder



Model Architecture

- 2D ViT encoder
- 3D ViT decoder

